

▼ Detect Covid-19 with Chest X-Ray using PyTorch

Welcome to detecting Covid-19 with Chest X-ray using PyTorch!

Student: Holger Espinola

Instructor: Amit Yadav

About this course

In 2 hour long guided project, we will use a ResNet-18 model and train it on a Covid-19 Radiography dataset. This dataset has nearly 3000 Chest X-ray scans which are categorized in 3 classes:

- Normal
- Viral pneumonia
- Covid-19

Our objective in this project is to create an image classification model that can predict Chest X-ray scans that belong to one of 3 classes with a reasonably high accuracy. Please note that this dataset, and the model that we train in the project, can not be used to diagnose Covid-19 or viral pneumonia. We are only using this data for educational purpose.

Before you attempt this project, you should be familiar with:

- Programming python
- Theoretical understanding of CNN
- Understanding optimization techniques such as gradient descent
- This is a hands on, practical project that focuses primarily on implementation, and not on the theory of CNN.

Before diving into the project, please take a look at the course objectives and structure:

Course Objectives

In this course, we are going to focus on the following learning objectives:

1. Create custom dataset and dataloader in PyTorch
2. Train a ResNet-18 model in PyTorch to perform Image Classification

By the end of this course, you will be able to create a CNN, and will be able to train it to classify Chest X-Ray scans with reasonably high accuracy.

Course Structure

This course is divided into 3 parts:

1. Course overview: This introductory reading material

2. Detecting Covid-19 with Chest X-Ray using PyTorch: This is the hands on project that we will work on in Rhyme.
3. Graded Quiz: This is the final assignment that you need to pass in order to finish the course successfully.

Project Structure

The hands on project on detecting Covid-19 with Chest X-Ray using PyTorch is divided into following tasks:

1. Task 1: Introduction
2. Task 2: Importing libraries
3. Task 3: Creating custom dataset
4. Task 4: Image transformations
5. Task 5: Prepare DataLoader
6. Task 6: Data Visualization
7. Task 7: Creating the model
8. Task 8: Training the model
9. Task 9: Final results

Meet the instructor

Amit Yadav is a Machine Learning Engineer with focus in creating Deep Learning based Computer Vision and Signal Processing products. He has led chat bot development at a large corporation in the past. Amit is one of the Machine Learning and Data Science instructors at Rhyme.

Earn a Certificate

After you have completed the detecting Covid-19 with Chest X-Ray using PyTorch hands-on project, you will be able to assess your knowledge using an ungraded assignment. Once you are comfortable with the concepts, take the final quiz, score higher than 80% to earn your certificate.

Learner Notebook

▼ Task 1. Introduction

```
1 #add files of google colab
2 from google.colab import drive
3 drive.mount('/content/drive')
```

Mounted at /content/drive

▼ Task 2. Importing Libraries

```

1 #import libraries
2 %matplotlib inline
3 import os
4 import shutil
5 import random
6 import torch
7 import torchvision
8 import numpy as np
9 from PIL import Image
10 import matplotlib.pyplot as plt
11
12 #make seed to random
13 torch.manual_seed(0)
14
15 #print version of pytorch
16 print('Using PyTorch version ', torch.__version__)
17

```

Using PyTorch version 1.7.0+cu101

Preparing training and test sets

```

1 !ls '/content/drive/My Drive/Pytorch Scholarship Challenge/Detect Covid with x-ray/COVID-19'
COVID-19          NORMAL.metadata.xlsx  'Viral Pneumonia.matadata.xlsx'
COVID-19.metadata.xlsx  README.md.txt
NORMAL            'Viral Pneumonia'

1 #define classes
2 class_names = ['normal', 'viral', 'covid']
3 root_dir = '/content/drive/My Drive/Pytorch Scholarship Challenge/Detect Covid with x-r
4 source_dirs = ['NORMAL', 'Viral Pneumonia', 'COVID-19']
5
6 if os.path.isdir(os.path.join(root_dir, source_dirs[1])):
7     os.mkdir(os.path.join(root_dir, 'test'))
8
9     #define names of images of training set
10    for i, d in enumerate(source_dirs):
11        os.rename(os.path.join(root_dir, d), os.path.join(root_dir, class_names[i]))
12
13    for c in class_names:
14        os.mkdir(os.path.join(root_dir, 'test', c))
15
16    #define 30 samples of images of test set
17    for c in class_names:
18        images = [x for x in os.listdir(os.path.join(root_dir, c)) if x.lower().endswith('.p
19        selected_images = random.sample(images, 30)

```

```

19     selected_images = random.sample(images, 50)
20
21     for image in selected_images:
22         source_path = os.path.join(root_dir, c, image)
23         target_path = os.path.join(root_dir, 'test', c, image)
24         shutil.move(source_path, target_path)
25

```

```
1 !ls '/content/drive/My Drive/Pytorch Scholarship Challenge/Detect Covid with x-ray/COVID'

```

```

covid          NORMAL.metadata.xlsx    viral
COVID-19.metadata.xlsx  README.md.txt      'Viral Pneumonia.matadata.xlsx'
normal         test

```

▼ Task 3. Creating custom dataset

```

1 class ChestXRayDataset(torch.utils.data.Dataset):
2     def __init__(self, image_dirs, transform):
3         def get_images(class_name):
4             images = [x for x in os.listdir(image_dirs[class_name]) if x.lower().endswith('.png')]
5             print(f'Found {len(images)} {class_name} examples')
6             return images
7
8         self.images = {}
9         self.class_names = ['normal', 'viral', 'covid']
10
11         for c in self.class_names:
12             self.images[c] = get_images(c)
13
14         self.image_dirs = image_dirs
15         self.transform = transform
16
17     def __len__(self):
18         return sum([len(self.images[c]) for c in self.class_names])
19
20     def __getitem__(self, index):
21         class_name = random.choice(self.class_names)
22         index = index % len(self.images[class_name])
23         image_name = self.images[class_name][index]
24         image_path = os.path.join(self.image_dirs[class_name], image_name)
25         image = Image.open(image_path).convert('RGB')
26         return self.transform(image), self.class_names.index(class_name)
27

```

▼ Task 4. Image transformation

```

1 #transformation to train set
2 train_transform = torchvision.transforms.Compose([
3     torchvision.transforms.Resize(size = (224, 224)),
4     torchvision.transforms.RandomHorizontalFlip(),
5     torchvision.transforms.ToTensor(),

```

```

6     torchvision.transforms.Normalize(mean = [0.485, 0.456, 0.406],
7                                       std = [0.229, 0.224, 0.225])
8 ])
9
10 #transformation to test set
11 test_transform = torchvision.transforms.Compose([
12     torchvision.transforms.Resize(size = (224, 224)),
13     torchvision.transforms.ToTensor(),
14     torchvision.transforms.Normalize(mean = [0.485, 0.456, 0.406],
15                                       std = [0.229, 0.224, 0.225])
16 ])

```

▼ Task 5. Prepare DataLoader

```
1 path_dir = '/content/drive/My Drive/Pytorch Scholarship Challenge/Detect Covid with x-r
```

```

1 train_dirs = {
2     'normal' : path_dir + 'COVID-19 Radiography Database/normal',
3     'viral' : path_dir + 'COVID-19 Radiography Database/viral',
4     'covid' : path_dir + 'COVID-19 Radiography Database/covid'
5 }
6
7 train_dataset = ChestXRayDataset(train_dirs, train_transform)

```

```

Found 1311 normal examples
Found 1315 viral examples
Found 189 covid examples

```

```

1 test_dirs = {
2     'normal' : path_dir + 'COVID-19 Radiography Database/test/normal',
3     'viral' : path_dir + 'COVID-19 Radiography Database/test/viral',
4     'covid' : path_dir + 'COVID-19 Radiography Database/test/covid'
5 }
6
7 test_dataset = ChestXRayDataset(test_dirs, test_transform)

```

```

Found 30 normal examples
Found 30 viral examples
Found 30 covid examples

```

```

1 #results of data loader
2 batch_size = 6
3
4 #dataloader of train
5 dl_train = torch.utils.data.DataLoader(train_dataset,
6                                         batch_size = batch_size,
7                                         shuffle = True)
8 #dataloader of test
9 dl_test = torch.utils.data.DataLoader(test_dataset,
10                                       batch_size = batch_size,
11                                       shuffle = True)
12 #print results

```

```

12 #print results
13 print(f'Num of training batches: {len(dl_train)}')
14 print(f'Num of testing batches: {len(dl_test)}')
15

```

Num of training batches: 470

Num of testing batches: 15

▼ Task 6. Data Visualization

```


1 class_names = train_dataset.class_names
2
3 def show_images(images, labels, preds):
4     plt.figure(figsize = (8, 4))
5
6     for i, image in enumerate(images):
7         plt.subplot(1, 6, i+1, xticks = [], yticks = [])
8         image = image.numpy().transpose((1, 2, 0))
9         mean = np.array([0.485, 0.456, 0.406])
10        std = np.array([0.229, 0.224, 0.225])
11        image = image * std + mean
12        image = np.clip(image, 0., 1.)
13        plt.imshow(image)
14
15        col = 'green' if preds[i] == labels[i] else 'red'
16        plt.xlabel(f'{class_names[int(labels[i].numpy())]}')
17        plt.ylabel(f'{class_names[int(preds[i].numpy())]}', color = col)
18
19    plt.tight_layout()
20    plt.show()
21

```

```

1 #show images of training
2 images, labels = next(iter(dl_train))
3 show_images(images, labels, labels)
4 print(images.shape)

```




```
torch.Size([6, 3, 224, 224])
```

```

1 #show images of testing
2 images, labels = next(iter(dl_test))
3 show_images(images, labels, labels)
4 print(images.shape)

```



▼ Task 7. Creating the Model

```
1 #load pre-trained resnet-18 model
2 resnet18 = torchvision.models.resnet18(pretrained = True)
3 print(resnet18)
```

Downloading: "<https://download.pytorch.org/models/resnet18-5c106cde.pth>" to /root/.c
100% 44.7M/44.7M [00:01<00:00, 24.3MB/s]

```
ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace=True)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
    (1): BasicBlock(
      (conv1): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer2): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(64, 128, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(64, 128, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer3): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (layer4): Sequential(
    (0): BasicBlock(
      (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (downsample): Sequential(
        (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
        (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      )
    )
    (1): BasicBlock(
      (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
      (relu): ReLU(inplace=True)
      (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False)
      (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    )
  )
  (fc): Linear(in_features=512, out_features=3)
  (loss_fn): CrossEntropyLoss()
  (optimizer): Adam(resnet18.parameters(), lr=3e-5)
```



```

        (conv1): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
        (bn1): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_
        (relu): ReLU(inplace=True)
        (conv2): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
        (bn2): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_
    )
)
(layer3): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(128, 256, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_
    (downsample): Sequential(
      (0): Conv2d(128, 256, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
    (bn1): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
    (bn2): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_
  )
)
(layer4): Sequential(
  (0): BasicBlock(
    (conv1): Conv2d(256, 512, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1),
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_
    (downsample): Sequential(
      (0): Conv2d(256, 512, kernel_size=(1, 1), stride=(2, 2), bias=False)
      (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_
    )
  )
  (1): BasicBlock(
    (conv1): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
    (bn1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_
    (relu): ReLU(inplace=True)
    (conv2): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
    (bn2): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_
  )
)
(avgpool): AdaptiveAvgPool2d(output_size=(1, 1))
(fc): Linear(in features=512, out features=3, bias=True)


```

```

1 #define function to show predictions
2
3 def show_preds():
4     resnet18.eval()
5     images, labels = next(iter(dl_test))
6     outputs = resnet18(images)
7     _, preds = torch.max(outputs, 1)
8     show_images(images, labels, preds)
9

```

```
1 show_preds()
```



▼ Task 8. Training the model

```
1 def training(epochs):
2     print('Starting training....')
3
4     for e in range(0, epochs):
5         print('=' * 20)
6         print(f'Starting epoch {e+1}/{epochs}')
7         print('=' * 20)
8
9         train_loss = 0.
10        resnet18.train()
11
12        for train_step, (images, labels) in enumerate(dl_train):
13            optimizer.zero_grad()
14            outputs = resnet18(images)
15            loss = loss_fn(outputs, labels)
16            loss.backward()
17            optimizer.step()
18            train_loss += loss.item()
19
20        if train_step % 20 == 0:
21            print('Evaluating step...', train_step)
22            val_loss = 0.
23            accuracy = 0.
24            resnet18.eval()
25
26            for val_step, (images, labels) in enumerate(dl_test):
27                outputs = resnet18(images)
28                loss = loss_fn(outputs, labels)
29                val_loss += loss.item()
30                _, preds = torch.max(outputs, 1)
31                accuracy += sum((preds == labels).numpy())
32            val_loss /= (val_step+1)
33            accuracy = accuracy/len(test_dataset)
34            print(f'Val loss: {val_loss:.4f} ==> Accuracy: {accuracy:.3f}')
35            show_preds()
36
37            resnet18.train()
38
39            if accuracy > 0.95:
40                print('Performance condition satisfied...')
41                return
```

```
42
43     train_loss /= (train_step + 1)
44     print(f'Training loss: {train_loss:.4f}')
45

1 training(epochs = 1)
```

Starting training....

=====

Starting epoch 1/1

=====

Evaluating step... 0


Val loss: 0.8886 ==> Accuracy: 0.544



▼ Task 9. Final Results



1 show_preds()



Evaluating step... 40

Project Finalized ... Congrats Holger Espinola ! Enjoy Coursera!




1

covid viral covid covid covid covid


Evaluating step... 60

Val loss: 0.2734 ==> Accuracy: 0.911



Evaluating step... 80

Val loss: 0.2073 ==> Accuracy: 0.922



Evaluating step... 100

Val loss: 0.1806 ==> Accuracy: 0.933

