


```
1 # set tf 1.x for colab
2 %tensorflow_version 1.x
```

Generating human faces with Adversarial Networks

 © research.nvidia.com_

This time we'll train a neural net to generate plausible human faces in all their subtlety: appearance, expression, accessories, etc. 'Cuz when us machines gonna take over Earth, there won't be any more faces left. We want to preserve this data for future iterations. Yikes...

Based on <https://github.com/Lasagne/Recipes/pull/94>.

▼ Running on Google Colab

```
1 ! shred -u setup_google_colab.py
2 ! wget https://raw.githubusercontent.com/hse-aml/intro-to-dl/master/setup_google_colab.py
3 import setup_google_colab
4 setup_google_colab.setup_week4()
```

```
--2021-02-02 23:16:25-- https://raw.githubusercontent.com/hse-aml/intro-to-dl/master/
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 151.101.0.133, 151.101.0.133, 151.101.0.133, 151.101.0.133
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|151.101.0.133|:443: connected.
HTTP request sent, awaiting response... 200 OK
Length: 3636 (3.6K) [text/plain]
Saving to: 'setup_google_colab.py'
```

```
setup_google_colab. 100%[=====>] 3.55K --.-KB/s in 0s
```

```
2021-02-02 23:16:26 (64.3 MB/s) - 'setup_google_colab.py' saved [3636/3636]
```

```
*****
lfw-deepfunneled.tgz
*****
lfw.tgz
*****
lfw_attributes.txt
```

```
1 import sys
2 sys.path.append("..")
3 import grading
4 import download_utils
5 import tqdm_utils
```

```
1 download_utils.link_week_4_resources()
```

```
1 import matplotlib.pyplot as plt
2 %matplotlib inline
3 import numpy as np
```

```

4 plt.rcParams.update({'axes.titlesize': 'small'})
5
6 from sklearn.datasets import load_digits
7 #The following line fetches you two datasets: images, usable for autoencoder training a
8 #Those attributes will be required for the final part of the assignment (applying smile
9 from lfw_dataset import load_lfw_dataset
10 data,attrs = load_lfw_dataset(dimx=36,dimy=36)
11
12 #preprocess faces
13 data = np.float32(data)/255.
14
15 IMG_SHAPE = data.shape[1:]

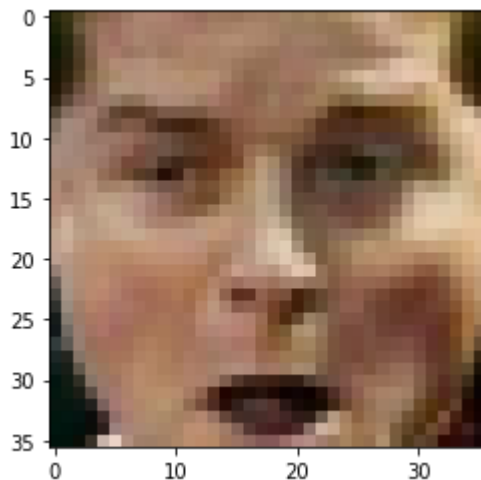
```

```


1 #print random image
2 plt.imshow(data[np.random.randint(data.shape[0])], cmap="gray", interpolation="none")

```

<matplotlib.image.AxesImage at 0x7f1782c3b978>



▼ Generative adversarial nets 101

 © torch.github.io_

Deep learning is simple, isn't it?

- build some network that generates the face (small image)
- make up a **measure of how good that face is**
- optimize with gradient descent :)

The only problem is: how can we engineers tell well-generated faces from bad? And i bet you we won't ask a designer for help.

If we can't tell good faces from bad, we delegate it to yet another neural network!

That makes the two of them:

- **Generator** - takes random noise for inspiration and tries to generate a face sample.

- Let's call him $G(z)$, where z is a gaussian noise.
- **Discriminator** - takes a face sample and tries to tell if it's great or fake.
 - Predicts the probability of input image being a **real face**
 - Let's call him $D(x)$, x being an image.
 - $D(x)$ is a prediction for real image and $D(G(z))$ is prediction for the face made by generator.

Before we dive into training them, let's construct the two networks.

```
1 import tensorflow as tf
2 from keras_utils import reset_tf_session
3 s = reset_tf_session()
4
5 import keras
6 from keras.models import Sequential
7 from keras import layers as L
```

Using TensorFlow backend.

```
WARNING:tensorflow:From /content/keras_utils.py:68: The name tf.get_default_session is deprecated. Please use tf.compat.v1.get_default_session instead.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:104: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:104: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:104: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.
WARNING:tensorflow:From /content/keras_utils.py:75: The name tf.ConfigProto is deprecated. Please use tf.compat.v1.ConfigProto instead.
WARNING:tensorflow:From /content/keras_utils.py:77: The name tf.InteractiveSession is deprecated. Please use tf.compat.v1.Session instead.
```

```
1 CODE_SIZE = 256
2
3 generator = Sequential()
4 generator.add(L.InputLayer([CODE_SIZE],name='noise'))
5 generator.add(L.Dense(10*8*8, activation='elu'))
6
7 generator.add(L.Reshape((8,8,10)))
8 generator.add(L.Deconv2D(64,kernel_size=(5,5),activation='elu'))
9 generator.add(L.Deconv2D(64,kernel_size=(5,5),activation='elu'))
10 generator.add(L.UpSampling2D(size=(2,2)))
11 generator.add(L.Deconv2D(32,kernel_size=3,activation='elu'))
12 generator.add(L.Deconv2D(32,kernel_size=3,activation='elu'))
13 generator.add(L.Deconv2D(32,kernel_size=3,activation='elu'))
14
15 generator.add(L.Conv2D(3,kernel_size=3,activation=None))
16
```

```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:104: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:104: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.
```

```
1 generator.summary()
```

Layer (type)	Output Shape	Param #
noise (InputLayer)	(None, 256)	0
dense_1 (Dense)	(None, 640)	164480
reshape_1 (Reshape)	(None, 8, 8, 10)	0
conv2d_transpose_1 (Conv2DTr	(None, 12, 12, 64)	16064
conv2d_transpose_2 (Conv2DTr	(None, 16, 16, 64)	102464
up_sampling2d_1 (UpSampling2	(None, 32, 32, 64)	0
conv2d_transpose_3 (Conv2DTr	(None, 34, 34, 32)	18464
conv2d_transpose_4 (Conv2DTr	(None, 36, 36, 32)	9248
conv2d_transpose_5 (Conv2DTr	(None, 38, 38, 32)	9248
conv2d_1 (Conv2D)	(None, 36, 36, 3)	867
Total params: 320,835		
Trainable params: 320,835		
Non-trainable params: 0		

```
1 assert generator.output_shape[1:] == IMG_SHAPE, "generator must output an image of shape"
```

▼ Discriminator

- Discriminator is your usual convolutional network with interlooping convolution and pooling layers
- The network does not include dropout/batchnorm to avoid learning complications.
- We also regularize the pre-output layer to prevent discriminator from being too certain.

```
1 IMG_SHAPE
```

```
(36, 36, 3)
```

```
1 discriminator = Sequential()
2
3 discriminator.add(L.InputLayer(IMG_SHAPE))
4
5 #<build discriminator body>
6 discriminator.add(L.Conv2D(filters = 16, kernel_size = (5, 5), activation = 'elu'))
7 discriminator.add(L.MaxPooling2D(pool_size = (2, 2)))
8 discriminator.add(L.Conv2D(filters = 32, kernel_size = (3, 3), padding = 'same',
9                               activation = 'elu'))
10 discriminator.add(L.Conv2D(filters = 32, kernel_size = (3, 3), padding = 'same',
11                              activation = 'elu'))
```

```

10 discriminator.add(L.MaxPooling2D(pool_size = (2, 2)))
11 discriminator.add(L.Conv2D(filters = 64, kernel_size = (3, 3), padding = 'same',
12                             activation = 'elu'))
13 discriminator.add(L.MaxPooling2D(pool_size = (2, 2)))
14 #
15 discriminator.add(L.Flatten())
16 discriminator.add(L.Dense(256,activation='tanh'))
17 discriminator.add(L.Dense(2,activation=tf.nn.log_softmax))
18

```

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorf

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorf
Instructions for updating:
keep_dims is deprecated, use keepdims instead


```
1 discriminator.summary()
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 36, 36, 3)	0
conv2d_2 (Conv2D)	(None, 32, 32, 16)	1216
max_pooling2d_1 (MaxPooling2	(None, 16, 16, 16)	0
conv2d_3 (Conv2D)	(None, 16, 16, 32)	4640
max_pooling2d_2 (MaxPooling2	(None, 8, 8, 32)	0
conv2d_4 (Conv2D)	(None, 8, 8, 64)	18496
max_pooling2d_3 (MaxPooling2	(None, 4, 4, 64)	0
flatten_1 (Flatten)	(None, 1024)	0
dense_2 (Dense)	(None, 256)	262400
dense_3 (Dense)	(None, 2)	514
Total params: 287,266		
Trainable params: 287,266		
Non-trainable params: 0		

▼ Training

We train the two networks concurrently:

- Train **discriminator** to better distinguish real data from **current** generator
- Train **generator** to make discriminator think generator is real
- Since discriminator is a differentiable neural network, we train both with gradient descent.

 © deeplearning4j.org

Training is done iteratively until discriminator is no longer able to find the difference (or until you run out of patience).

Tricks:

- Regularize discriminator output weights to prevent explosion
- Train generator with **adam** to speed up training. Discriminator trains with SGD to avoid problems with momentum.
- More: <https://github.com/soumith/ganhacks>

```

1 noise = tf.placeholder('float32',[None,CODE_SIZE])
2 real_data = tf.placeholder('float32',[None,]+list(IMG_SHAPE))
3
4 logp_real = discriminator(real_data)
5
6 #generated_data = <gen(noise)>
7 generated_data = generator(noise)
8
9 #logp_gen = <log P(real | gen(noise))>
10 logp_gen = discriminator(generated_data)
11

```

```

1 #####
2 #discriminator training#
3 #####
4
5 d_loss = -tf.reduce_mean(logp_real[:,1] + logp_gen[:,0])
6
7 #regularize
8 d_loss += tf.reduce_mean(discriminator.layers[-1].kernel**2)
9
10 #optimize
11 disc_optimizer = tf.train.GradientDescentOptimizer(1e-3).minimize(d_loss,var_list=disc

```

WARNING:tensorflow:From /tensorflow-1.15.2/python3.6/tensorflow_core/python/ops/math_ops.py:1226: tf.nn.conv2d is deprecated and will be removed in a future version. Instructions for updating:
Use tf.nn.conv2d in 2.0, which has the same broadcast rule as np.where

```

1 #####
2 ###generator training###
3 #####
4
5 #g_loss = <generator loss>
6 g_loss = -tf.reduce_mean(logp_gen[:, 1])
7
8 gen_optimizer = tf.train.AdamOptimizer(1e-4).minimize(g_loss,var_list=generator.trainab
9
10

```

```
1 s.run(tf.global_variables_initializer())
```

▼ Auxiliary functions

Here we define a few helper functions that draw current data distributions and sample training batches.

```
1 def sample_noise_batch(bsize):
2     return np.random.normal(size=(bsize, CODE_SIZE)).astype('float32')
3
4 def sample_data_batch(bsize):
5     idxs = np.random.choice(np.arange(data.shape[0]), size=bsize)
6     return data[idxs]
7
8 def sample_images(nrow,ncol, sharp=False):
9     images = generator.predict(sample_noise_batch(bsize=nrow*ncol))
10    if np.var(images)!=0:
11        images = images.clip(np.min(data),np.max(data))
12    for i in range(nrow*ncol):
13        plt.subplot(nrow,ncol,i+1)
14        if sharp:
15            plt.imshow(images[i].reshape(IMG_SHAPE),cmap="gray", interpolation="none")
16        else:
17            plt.imshow(images[i].reshape(IMG_SHAPE),cmap="gray")
18    plt.show()
19
20 def sample_probas(bsize):
21     plt.title('Generated vs real data')
22     plt.hist(np.exp(discriminator.predict(sample_data_batch(bsize)))[:,1],
23             label='D(x)', alpha=0.5,range=[0,1])
24     plt.hist(np.exp(discriminator.predict(generator.predict(sample_noise_batch(bsize))))
25             label='D(G(z))',alpha=0.5,range=[0,1])
26     plt.legend(loc='best')
27     plt.show()
```

▼ Training

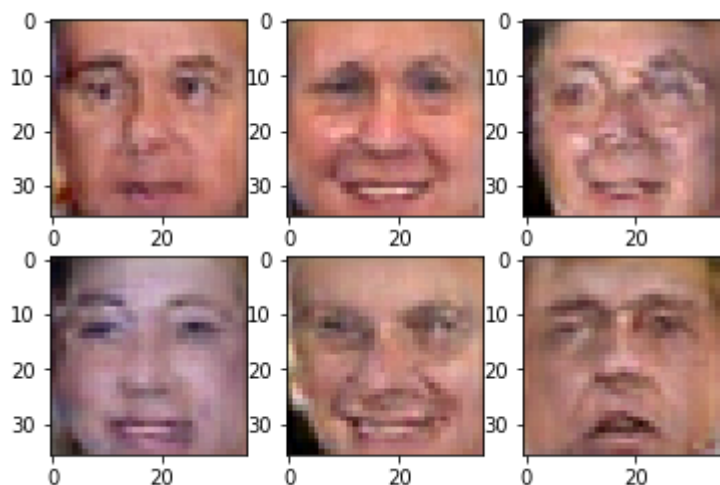
Main loop. We just train generator and discriminator in a loop and plot results once every N iterations.

```
1 from IPython import display
2
3 for epoch in tqdm_utils.tqdm_notebook_failsafe(range(50000)):
4
5     feed_dict = {
6         real_data:sample_data_batch(100),
7         noise:sample_noise_batch(100)
8     }
9
10    for i in range(5):
11        s.run(disc_optimizer,feed_dict)
```

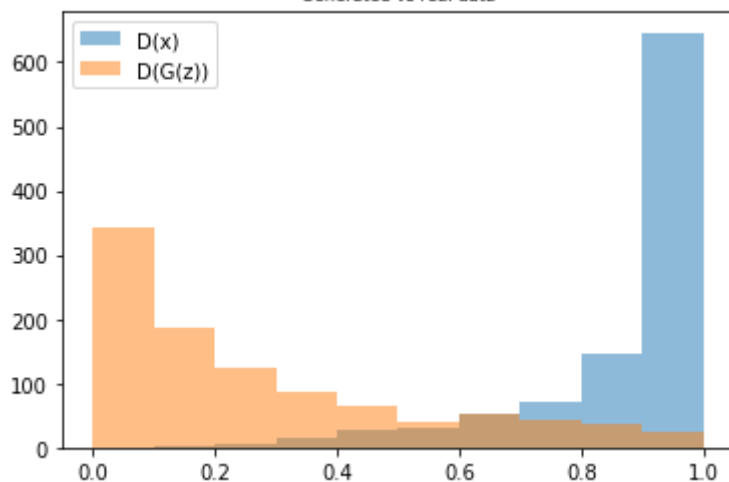
```

12
13     s.run(gen_optimizer, feed_dict)
14
15     if epoch % 100 == 0:
16         display.clear_output(wait=True)
17         sample_images(2, 3, True)
18         sample_probas(1000)
19

```



Generated vs real data



*

```

1 from submit_honor import submit_honor
2 #submit_honor((generator, discriminator), <YOUR_EMAIL>, <YOUR_TOKEN>)
3 submit_honor((generator, discriminator), 'knowtech94@gmail.com', '1lRyCzY9UQ027NpB')


```

Submitted to Coursera platform. See results on assignment page!

```

1 #The network was trained for about 15k iterations.
2 #Training for longer yields MUCH better results
3 plt.figure(figsize=[16,24])
4 sample_images(16,8)

```

1