

week6_final_project_image_captioning_clean

July 28, 2021

1 Image Captioning Final Project

In this final project you will define and train an image-to-caption model, that can produce descriptions for real world images!

Model architecture: CNN encoder and RNN decoder. (<https://research.googleblog.com/2014/11/a-picture-is-worth-thousand-coherent.html>)

2 Import stuff

```
In [1]: import sys
        sys.path.append("..")
        import grading
        import download_utils

In [2]: download_utils.link_all_keras_resources()

In [3]: import tensorflow as tf
        from tensorflow.contrib import keras
        import numpy as np
        %matplotlib inline
        import matplotlib.pyplot as plt
        L = keras.layers
        K = keras.backend
        import utils
        import time
        import zipfile
        import json
        from collections import defaultdict
        import re
        import random
        from random import choice
        import grading_utils
        import os
        from keras_utils import reset_tf_session
        import tqdm_utils
```

Using TensorFlow backend.

3 Fill in your Coursera token and email

To successfully submit your answers to our grader, please fill in your Coursera submission token and email

```
In [4]: grader = grading.Grader(assignment_key="NEDBg6CgEee8nQ6uE8a70A",
                                all_parts=["19Wpv", "uJh73", "yiJkt", "rbpnH", "E20IL", "YJR7z"])

In [34]: # token expires every 30 min
          COURSERA_TOKEN = 'ShgxVTpxTGEgIwcd'      ### YOUR TOKEN HERE
          COURSERA_EMAIL = 'knowtech94@gmail.com'    ### YOUR EMAIL HERE
```

4 Download data

Takes 10 hours and 20 GB. We've downloaded necessary files for you.

Relevant links (just in case): - train images <http://msvocds.blob.core.windows.net/coco2014/train2014.zip>
- validation images <http://msvocds.blob.core.windows.net/coco2014/val2014.zip> - captions for both train and validation http://msvocds.blob.core.windows.net/annotations-1-0-3/captions_train-val2014.zip

```
In [8]: # we downloaded them for you, just link them here
        download_utils.link_week_6_resources()
```

5 Extract image features

We will use pre-trained InceptionV3 model for CNN encoder (<https://research.googleblog.com/2016/03/train-your-own-image-classifier-with.html>) and extract its last hidden layer as an embedding:

```
In [9]: IMG_SIZE = 299

In [10]: # we take the last hidden layer of IncetionV3 as an image embedding
         def get_cnn_encoder():
             K.set_learning_phase(False)
             model = keras.applications.InceptionV3(include_top=False)
             preprocess_for_model = keras.applications.inception_v3.preprocess_input

             model = keras.models.Model(model.inputs, keras.layers.GlobalAveragePooling2D()(model.outputs))
             return model, preprocess_for_model
```

Features extraction takes too much time on CPU: - Takes 16 minutes on GPU. - 25x slower (InceptionV3) on CPU and takes 7 hours. - 10x slower (MobileNet) on CPU and takes 3 hours.

So we've done it for you with the following code:

```
# load pre-trained model
reset_tf_session()
encoder, preprocess_for_model = get_cnn_encoder()
```

```

# extract train features
train_img_embeds, train_img_fns = utils.apply_model(
    "train2014.zip", encoder, preprocess_for_model, input_shape=(IMG_SIZE, IMG_SIZE))
utils.save_pickle(train_img_embeds, "train_img_embeds.pickle")
utils.save_pickle(train_img_fns, "train_img_fns.pickle")

# extract validation features
val_img_embeds, val_img_fns = utils.apply_model(
    "val2014.zip", encoder, preprocess_for_model, input_shape=(IMG_SIZE, IMG_SIZE))
utils.save_pickle(val_img_embeds, "val_img_embeds.pickle")
utils.save_pickle(val_img_fns, "val_img_fns.pickle")

# sample images for learners
def sample_zip(fn_in, fn_out, rate=0.01, seed=42):
    np.random.seed(seed)
    with zipfile.ZipFile(fn_in) as fin, zipfile.ZipFile(fn_out, "w") as fout:
        sampled = filter(lambda _: np.random.rand() < rate, fin.filelist)
        for zInfo in sampled:
            fout.writestr(zInfo, fin.read(zInfo))

sample_zip("train2014.zip", "train2014_sample.zip")
sample_zip("val2014.zip", "val2014_sample.zip")

```

```

In [11]: # load prepared embeddings
train_img_embeds = utils.read_pickle("train_img_embeds.pickle")
train_img_fns = utils.read_pickle("train_img_fns.pickle")
val_img_embeds = utils.read_pickle("val_img_embeds.pickle")
val_img_fns = utils.read_pickle("val_img_fns.pickle")
# check shapes
print(train_img_embeds.shape, len(train_img_fns))
print(val_img_embeds.shape, len(val_img_fns))

```

```

(82783, 2048) 82783
(40504, 2048) 40504

```

```

In [12]: # check prepared samples of images
list(filter(lambda x: x.endswith("_sample.zip"), os.listdir(".")))

```

```

Out[12]: ['val2014_sample.zip', 'train2014_sample.zip']

```

6 Extract captions for images

```

In [13]: # extract captions from zip
def get_captions_for_fns(fns, zip_fn, zip_json_path):
    zf = zipfile.ZipFile(zip_fn)
    j = json.loads(zf.read(zip_json_path).decode("utf8"))
    id_to_fn = {img["id"]: img["file_name"] for img in j["images"]}

```

```

fn_to_caps = defaultdict(list)
for cap in j['annotations']:
    fn_to_caps[id_to_fn[cap['image_id']]].append(cap['caption'])
fn_to_caps = dict(fn_to_caps)
return list(map(lambda x: fn_to_caps[x], fns))

train_captions = get_captions_for_fns(train_img_fns, "captions_train-val2014.zip",
                                     "annotations/captions_train2014.json")

val_captions = get_captions_for_fns(val_img_fns, "captions_train-val2014.zip",
                                    "annotations/captions_val2014.json")

# check shape
print(len(train_img_fns), len(train_captions))
print(len(val_img_fns), len(val_captions))

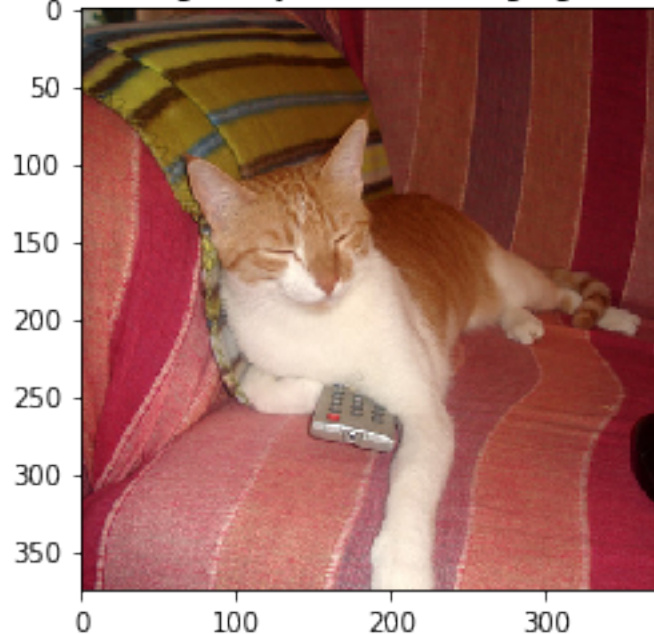
82783 82783
40504 40504

In [14]: # look at training example (each has 5 captions)
def show_trainig_example(train_img_fns, train_captions, example_idx=0):
    """
    You can change example_idx and see different images
    """
    zf = zipfile.ZipFile("train2014_sample.zip")
    captions_by_file = dict(zip(train_img_fns, train_captions))
    all_files = set(train_img_fns)
    found_files = list(filter(lambda x: x.filename.rsplit("/")[-1] in all_files, zf.files))
    example = found_files[example_idx]
    img = utils.decode_image_from_buf(zf.read(example))
    plt.imshow(utils.image_center_crop(img))
    plt.title("\n".join(captions_by_file[example.filename.rsplit("/")[-1]]))
    plt.show()

show_trainig_example(train_img_fns, train_captions, example_idx=142)

```

A cat sitting on a pink stripped couch
 An orange and white cat sitting in a striped chair.
 An orange and white cat sleeping on a remote
 Brown and white cat sleeping on couch while lying on remote.
 A cat closing its eyes while lounging on a chair.



7 Prepare captions for training

```
In [15]: # preview captions data
         train_captions[:2]
```

```
Out[15]: [['A long dirt road going through a forest.',
           'A SCENE OF WATER AND A PATH WAY',
           'A sandy path surrounded by trees leads to a beach.',
           'Ocean view through a dirt road surrounded by a forested area. ',
           'dirt path leading beneath barren trees to open plains'],
          ['A group of zebra standing next to each other.',
           'This is an image of of zebras drinking',
           'ZEBRAS AND BIRDS SHARING THE SAME WATERING HOLE',
           'Zebras that are bent over and drinking water together.',
           'a number of zebras drinking water near one another']]
```

```
In [16]: # special tokens
         PAD = "#PAD#"
         UNK = "#UNK#"
         START = "#START#"
```

```

END = "#END#"

# split sentence into tokens (split into lowercased words)
def split_sentence(sentence):
    return list(filter(lambda x: len(x) > 0, re.split('\W+', sentence.lower())))

def generate_vocabulary(train_captions):
    """
    Return {token: index} for all train tokens (words) that occur 5 times or more,
    `index` should be from 0 to N, where N is a number of unique tokens in the results.
    Use `split_sentence` function to split sentence into tokens.
    Also, add PAD (for batch padding), UNK (unknown, out of vocabulary),
    START (start of sentence) and END (end of sentence) tokens into the vocabulary.
    """
    #vocab = ### YOUR CODE HERE ###

    words = [words for captions in train_captions for caption in captions for words in caption]
    words = np.reshape(words, -1)
    unique_words = np.unique(words)
    word_count = {word: 0 for word in unique_words}
    for word in unique_words:
        word_count[word] = word_count[word] + 1

    vocab = [w for w, c in word_count.items() if c >= 5]
    vocab.append(PAD)
    vocab.append(UNK)
    vocab.append(START)
    vocab.append(END)

    return {pal:count for count, pal in enumerate(sorted(vocab))}

def caption_tokens_to_indices(captions, vocab):
    """
    `captions` argument is an array of arrays:
    [
        [
            "image1 caption1",
            "image1 caption2",
            ...
        ],
        [
            "image2 caption1",
            "image2 caption2",
            ...
        ],
        ...
    ]
    Use `split_sentence` function to split sentence into tokens.
    """

```

Replace all tokens with vocabulary indices, use UNK for unknown words (out of vocab)
Add START and END tokens to start and end of each sentence respectively.
For the example above you should produce the following:

```
[
    [
        [vocab[START], vocab["image1"], vocab["caption1"], vocab[END]],
        [vocab[START], vocab["image1"], vocab["caption2"], vocab[END]],
        ...
    ],
    ...
]
"""
#res = ### YOUR CODE HERE ###

vocab_words = lambda words: [vocab[START]] + [vocab[word] if word in vocab else vocab[UNK]]
total_words = lambda captions: [vocab_words(split_sentence(caption.lower())) for caption in captions]
response = [total_words(caption_list) for caption_list in captions]

return response
```

```
In [17]: # prepare vocabulary
vocab = generate_vocabulary(train_captions)
vocab_inverse = {idx: w for w, idx in vocab.items()}
print(len(vocab))
```

8769

```
In [18]: # replace tokens with indices
train_captions_indexed = caption_tokens_to_indices(train_captions, vocab)
val_captions_indexed = caption_tokens_to_indices(val_captions, vocab)
```

Captions have different length, but we need to batch them, that's why we will add PAD tokens so that all sentences have an equal length.

We will crunch LSTM through all the tokens, but we will ignore padding tokens during loss calculation.

```
In [19]: # we will use this during training
def batch_captions_to_matrix(batch_captions, pad_idx, max_len=None):
    """
    `batch_captions` is an array of arrays:
    [
        [vocab[START], ..., vocab[END]],
        [vocab[START], ..., vocab[END]],
        ...
    ]
    Put vocabulary indexed captions into np.array of shape (len(batch_captions), columns)
    where "columns" is max(map(len, batch_captions)) when max_len is None
    and "columns" = min(max_len, max(map(len, batch_captions))) otherwise.
```

```

Add padding with pad_idx where necessary.
Input example: [[1, 2, 3], [4, 5]]
Output example: np.array([[1, 2, 3], [4, 5, pad_idx]]) if max_len=None
Output example: np.array([[1, 2], [4, 5]]) if max_len=2
Output example: np.array([[1, 2, 3], [4, 5, pad_idx]]) if max_len=100
Try to use numpy, we need this function to be fast!
"""

#matrix = ###YOUR CODE HERE###

max_len = max_len or max(map(len, batch_captions))
max_len = min(max_len, max(map(len, batch_captions)))
matrix = np.empty((len(batch_captions), max_len))
matrix.fill(pad_idx)
for i in range(len(batch_captions)):
    line_ix = list(batch_captions[i])[:max_len]
    matrix[i,:len(line_ix)] = line_ix

return matrix

```

```

In [20]: ## GRADED PART, DO NOT CHANGE!
# Vocabulary creation
grader.set_answer("19Wpv", grading_utils.test_vocab(vocab, PAD, UNK, START, END))
# Captions indexing
grader.set_answer("uJh73", grading_utils.test_captions_indexing(train_captions_indexed,
# Captions batching
grader.set_answer("yiJkt", grading_utils.test_captions_batching(batch_captions_to_matrix,
In [21]: # you can make submission with answers so far to check yourself at this stage
grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)

```

Submitted to Coursera platform. See results on assignment page!

```

In [22]: # make sure you use correct argument in caption_tokens_to_indices
assert len(caption_tokens_to_indices(train_captions[:10], vocab)) == 10
assert len(caption_tokens_to_indices(train_captions[:5], vocab)) == 5

```

8 Training

8.1 Define architecture

Since our problem is to generate image captions, RNN text generator should be conditioned on image. The idea is to use image features as an initial state for RNN instead of zeros.

Remember that you should transform image feature vector to RNN hidden state size by fully-connected layer and then pass it to RNN.

During training we will feed ground truth tokens into the lstm to get predictions of next tokens.

Notice that we don't need to feed last token (END) as input (<http://cs.stanford.edu/people/karpathy/>):


```
In [23]: IMG_EMBED_SIZE = train_img_embeds.shape[1]
        IMG_EMBED_BOTTLENECK = 120
        WORD_EMBED_SIZE = 100
        LSTM_UNITS = 300
        LOGIT_BOTTLENECK = 120
        pad_idx = vocab[PAD]

In [24]: # remember to reset your graph if you want to start building it from scratch!
        s = reset_tf_session()
        tf.set_random_seed(42)
```

Here we define decoder graph.

We use Keras layers where possible because we can use them in functional style with weights reuse like this:

```
dense_layer = L.Dense(42, input_shape=(None, 100) activation='relu')
a = tf.placeholder('float32', [None, 100])
b = tf.placeholder('float32', [None, 100])
dense_layer(a) # that's how we applied dense layer!
dense_layer(b) # and again
```

Here's a figure to help you with flattening in decoder:

```
In [25]: class decoder:
        # [batch_size, IMG_EMBED_SIZE] of CNN image features
        img_embeds = tf.placeholder('float32', [None, IMG_EMBED_SIZE])
        # [batch_size, time steps] of word ids
        sentences = tf.placeholder('int32', [None, None])

        # we use bottleneck here to reduce the number of parameters
        # image embedding -> bottleneck
        img_embed_to_bottleneck = L.Dense(IMG_EMBED_BOTTLENECK,
                                           input_shape=(None, IMG_EMBED_SIZE),
                                           activation='elu')
        # image embedding bottleneck -> lstm initial state
        img_embed_bottleneck_to_h0 = L.Dense(LSTM_UNITS,
                                              input_shape=(None, IMG_EMBED_BOTTLENECK),
                                              activation='elu')

        # word -> embedding
        word_embed = L.Embedding(len(vocab), WORD_EMBED_SIZE)
        # lstm cell (from tensorflow)
        lstm = tf.nn.rnn_cell.LSTMCell(LSTM_UNITS)

        # we use bottleneck here to reduce model complexity
        # lstm output -> logits bottleneck
        token_logits_bottleneck = L.Dense(LOGIT_BOTTLENECK,
                                           input_shape=(None, LSTM_UNITS),
                                           activation="elu")
        # logits bottleneck -> logits for next token prediction
```

```

token_logits = L.Dense(len(vocab),
                        input_shape=(None, LOGIT_BOTTLENECK))

# initial lstm cell state of shape (None, LSTM_UNITS),
# we need to condition it on `img_embeds` placeholder.
#c0 = h0 = ### YOUR CODE HERE ###
c0 = h0 = img_embed_bottleneck_to_h0(img_embed_to_bottleneck(img_embeds))

# embed all tokens but the last for lstm input,
# remember that L.Embedding is callable,
# use `sentences` placeholder as input.
#word_embeds = ### YOUR CODE HERE ###
word_embeds = word_embed(sentences[:, :-1])

# during training we use ground truth tokens `word_embeds` as context for next tokens
# that means that we know all the inputs for our lstm and can get
# all the hidden states with one tensorflow operation (tf.nn.dynamic_rnn).
# `hidden_states` has a shape of [batch_size, time steps, LSTM_UNITS].
hidden_states, _ = tf.nn.dynamic_rnn(lstm, word_embeds,
                                     initial_state=tf.nn.rnn_cell.LSTMStateTuple(c0, h0))

# now we need to calculate token logits for all the hidden states

# first, we reshape `hidden_states` to [-1, LSTM_UNITS]
#flat_hidden_states = ### YOUR CODE HERE ###
flat_hidden_states = tf.reshape(hidden_states, [-1, LSTM_UNITS])

# then, we calculate logits for next tokens using `token_logits_bottleneck` and `token_logits`
#flat_token_logits = ### YOUR CODE HERE ###
flat_token_logits = token_logits(token_logits_bottleneck(flat_hidden_states))

# then, we flatten the ground truth token ids.
# remember, that we predict next tokens for each time step,
# use `sentences` placeholder.
#flat_ground_truth = ### YOUR CODE HERE ###
flat_ground_truth = tf.reshape(sentences[:, 1:], [-1])

# we need to know where we have real tokens (not padding) in `flat_ground_truth`,
# we don't want to propagate the loss for padded output tokens,
# fill `flat_loss_mask` with 1.0 for real tokens (not pad_idx) and 0.0 otherwise.
#flat_loss_mask = ### YOUR CODE HERE ###
flat_loss_mask = tf.not_equal(flat_ground_truth, pad_idx)

# compute cross-entropy between `flat_ground_truth` and `flat_token_logits` predictions
xent = tf.nn.sparse_softmax_cross_entropy_with_logits(
    labels=flat_ground_truth,
    logits=flat_token_logits
)

```

```

# compute average `xent` over tokens with nonzero `flat_loss_mask`.
# we don't want to account misclassification of PAD tokens, because that doesn't make sense
# we have PAD tokens for batching purposes only!
#loss = ### YOUR CODE HERE ###
loss = tf.reduce_mean(tf.boolean_mask(xent, flat_loss_mask))

```

```

In [26]: # define optimizer operation to minimize the loss
optimizer = tf.train.AdamOptimizer(learning_rate=0.001)
train_step = optimizer.minimize(decoder.loss)

# will be used to save/load network weights.
# you need to reset your default graph and define it in the same way to be able to load
saver = tf.train.Saver()

# initialize all variables
s.run(tf.global_variables_initializer())

```

```

/opt/conda/lib/python3.6/site-packages/tensorflow/python/ops/gradients_impl.py:93: UserWarning:
"Converting sparse IndexedSlices to a dense Tensor of unknown shape. "

```

```

In [27]: ## GRADED PART, DO NOT CHANGE!
# Decoder shapes test
grader.set_answer("rbpnH", grading_utils.test_decoder_shapes(decoder, IMG_EMBED_SIZE, vocab_size))
# Decoder random loss test
grader.set_answer("E20IL", grading_utils.test_random_decoder_loss(decoder, IMG_EMBED_SIZE, vocab_size))

```

```

In [28]: # you can make submission with answers so far to check yourself at this stage
grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)

```

Submitted to Coursera platform. See results on assignment page!

8.2 Training loop

Evaluate train and validation metrics through training and log them. Ensure that loss decreases.

```

In [29]: train_captions_indexed = np.array(train_captions_indexed)
val_captions_indexed = np.array(val_captions_indexed)

In [30]: # generate batch via random sampling of images and captions for them,
# we use `max_len` parameter to control the length of the captions (truncating long captions)
def generate_batch(images_embeddings, indexed_captions, batch_size, max_len=None):
    """
    `images_embeddings` is a np.array of shape [number of images, IMG_EMBED_SIZE].
    `indexed_captions` holds 5 vocabulary indexed captions for each image:
    [

```

```

        [
            [vocab[START], vocab["image1"], vocab["caption1"], vocab[END]],
            [vocab[START], vocab["image1"], vocab["caption2"], vocab[END]],
            ...
        ],
        ...
    ]

    Generate a random batch of size `batch_size`.
    Take random images and choose one random caption for each image.
    Remember to use `batch_captions_to_matrix` for padding and respect `max_len` parameter.
    Return feed dict {decoder.img_embeds: ..., decoder.sentences: ...}.
    """
    #batch_image_embeddings = ### YOUR CODE HERE ###
    batch_start_idx = np.random.randint(0, high = len(images_embeddings) - batch_size)
    idxs = np.random.permutation(range(len(images_embeddings)))[batch_start_idx: batch_start_idx + batch_size]
    batch_image_embeddings = np.take(images_embeddings, idxs, axis = 0)

    #batch_captions_matrix = ### YOUR CODE HERE ###
    caption_inner_idxs = np.random.choice(range(5), batch_size)
    batch_captions_matrix = [indexed_captions[m][n] for m, n in zip(idxs, caption_inner_idxs)]
    batch_captions_matrix = batch_captions_to_matrix(batch_captions_matrix, pad_idx, max_len)

    return {decoder.img_embeds: batch_image_embeddings,
            decoder.sentences: batch_captions_matrix}

```

```

In [32]: batch_size = 64
        n_epochs = 10
        n_batches_per_epoch = 1000
        n_validation_batches = 100 # how many batches are used for validation after each epoch

```

```

In [30]: # you can load trained weights here
        # you can load "weights_{epoch}" and continue training
        # uncomment the next line if you need to load weights
        # saver.restore(s, os.path.abspath("weights"))

```

Look at the training and validation loss, they should be decreasing!

```

In [33]: # actual training loop
        MAX_LEN = 20 # truncate long captions to speed up training

        # to make training reproducible
        np.random.seed(42)
        random.seed(42)

        for epoch in range(n_epochs):

            train_loss = 0
            pbar = tqdm_utils.tqdm_notebook_failsafe(range(n_batches_per_epoch))
            counter = 0

```

```

for _ in pbar:
    train_loss += s.run([decoder.loss, train_step],
                        generate_batch(train_img_embeds,
                                      train_captions_indexed,
                                      batch_size,
                                      MAX_LEN))[0]

    counter += 1
    pbar.set_description("Training loss: %f" % (train_loss / counter))

train_loss /= n_batches_per_epoch

val_loss = 0
for _ in range(n_validation_batches):
    val_loss += s.run(decoder.loss, generate_batch(val_img_embeds,
                                                  val_captions_indexed,
                                                  batch_size,
                                                  MAX_LEN))

val_loss /= n_validation_batches

print('Epoch: {}, train loss: {}, val loss: {}'.format(epoch, train_loss, val_loss))

# save weights after finishing epoch
saver.save(s, os.path.abspath("weights_{}".format(epoch)))

print("Finished!")

```

A Jupyter Widget

Epoch: 0, train loss: 4.291811956882476, val loss: 3.625504195690155

A Jupyter Widget

Epoch: 1, train loss: 3.3434647524356844, val loss: 3.166829035282135

A Jupyter Widget

Epoch: 2, train loss: 3.0545015444755554, val loss: 2.9616402840614318

A Jupyter Widget

Epoch: 3, train loss: 2.903541965007782, val loss: 2.8957045578956606

A Jupyter Widget

Epoch: 4, train loss: 2.805233480453491, val loss: 2.825015242099762

A Jupyter Widget

Epoch: 5, train loss: 2.7360985076427458, val loss: 2.7711693978309633

A Jupyter Widget

Epoch: 6, train loss: 2.692197933912277, val loss: 2.7144248127937316

A Jupyter Widget

Epoch: 7, train loss: 2.6467096557617187, val loss: 2.6921755361557005

A Jupyter Widget

Epoch: 8, train loss: 2.6048688378334046, val loss: 2.674499177932739

A Jupyter Widget

Epoch: 9, train loss: 2.579082692861557, val loss: 2.6435467195510864
Finished!

```
In [35]: ## GRADED PART, DO NOT CHANGE!
         # Validation loss
         grader.set_answer("YJR7z", grading_utils.test_validation_loss(
             decoder, s, generate_batch, val_img_embeds, val_captions_indexed))
```

A Jupyter Widget

```
In [36]: # you can make submission with answers so far to check yourself at this stage
         grader.submit(COURSERA_EMAIL, COURSERA_TOKEN)
```

Submitted to Coursera platform. See results on assignment page!

```
In [37]: # check that it's learnt something, outputs accuracy of next word prediction (should be
         from sklearn.metrics import accuracy_score, log_loss
```

```
def decode_sentence(sentence_indices):
    return " ".join(list(map(vocab_inverse.get, sentence_indices)))

def check_after_training(n_examples):
    fd = generate_batch(train_img_embeds, train_captions_indexed, batch_size)
    logits = decoder.flat_token_logits.eval(fd)
    truth = decoder.flat_ground_truth.eval(fd)
    mask = decoder.flat_loss_mask.eval(fd).astype(bool)
    print("Loss:", decoder.loss.eval(fd))
    print("Accuracy:", accuracy_score(logits.argmax(axis=1)[mask], truth[mask]))
    for example_idx in range(n_examples):
        print("Example", example_idx)
        print("Predicted:", decode_sentence(logits.argmax(axis=1).reshape((batch_size,
        print("Truth:", decode_sentence(truth.reshape((batch_size, -1))[example_idx]))
        print("")

    check_after_training(3)
```

Loss: 2.74344

Accuracy: 0.444600280505

Example 0

Predicted: a woman eating and a bowl of a of a pizza #END# food #END# #END# #END# #END# #END#

Truth: a kid drink from a glass in front of a food of plate #END# #PAD# #PAD# #PAD# #PAD#

Example 1

Predicted: a toilets of a on a store on a middle #END# #END# #END# #END# #END# #END# #END# #END#

Truth: some bottles of cleaner in a spot in the wall #END# #PAD# #PAD# #PAD# #PAD# #PAD# #PAD# #PAD#

Example 2

Predicted: a bananas and on a outdoor plate box #END# a table plate #END# #END# #END# #END# #END#

Truth: sliced bananas sit on an open faced sandwich on a paper plate #END# #PAD# #PAD# #PAD# #PAD#

```
In [38]: # save graph weights to file!
saver.save(s, os.path.abspath("weights"))
```

```
Out[38]: '/home/jovyan/work/week6/weights'
```

9 Applying model

Here we construct a graph for our final model.

It will work as follows: - take an image as an input and embed it - condition lstm on that embedding - predict the next token given a START input token - use predicted token as an input at next time step - iterate until you predict an END token

```
In [39]: class final_model:
    # CNN encoder
    encoder, preprocess_for_model = get_cnn_encoder()
    saver.restore(s, os.path.abspath("weights")) # keras applications corrupt our graph

    # containers for current lstm state
    lstm_c = tf.Variable(tf.zeros([1, LSTM_UNITS]), name="cell")
    lstm_h = tf.Variable(tf.zeros([1, LSTM_UNITS]), name="hidden")

    # input images
    input_images = tf.placeholder('float32', [1, IMG_SIZE, IMG_SIZE, 3], name='images')

    # get image embeddings
    img_embeds = encoder(input_images)

    # initialize lstm state conditioned on image
    init_c = init_h = decoder.img_embed_bottleneck_to_h0(decoder.img_embed_to_bottleneck(img_embeds))
    init_lstm = tf.assign(lstm_c, init_c), tf.assign(lstm_h, init_h)

    # current word index
    current_word = tf.placeholder('int32', [1], name='current_input')

    # embedding for current word
    word_embed = decoder.word_embed(current_word)

    # apply lstm cell, get new lstm states
    new_c, new_h = decoder.lstm(word_embed, tf.nn.rnn_cell.LSTMStateTuple(lstm_c, lstm_h))

    # compute logits for next token
    new_logits = decoder.token_logits(decoder.token_logits_bottleneck(new_h))
    # compute probabilities for next token
    new_probs = tf.nn.softmax(new_logits)

    # `one_step` outputs probabilities of next token and updates lstm hidden state
    one_step = new_probs, tf.assign(lstm_c, new_c), tf.assign(lstm_h, new_h)
```



```
INFO:tensorflow:Restoring parameters from /home/jovyan/work/week6/weights
```

```
In [40]: # look at how temperature works for probability distributions
# for high temperature we have more uniform distribution
_ = np.array([0.5, 0.4, 0.1])
for t in [0.01, 0.1, 1, 10, 100]:
    print(" ".join(map(str, _*(1/t) / np.sum(_*(1/t)))), "with temperature", t)
```

```
0.999999999796 2.03703597592e-10 1.26765059997e-70 with temperature 0.01
0.903037043325 0.0969628642039 9.24709932365e-08 with temperature 0.1
0.5 0.4 0.1 with temperature 1
0.353447726392 0.345648113606 0.300904160002 with temperature 10
0.335367280481 0.334619764349 0.33001295517 with temperature 100
```

```
In [41]: # this is an actual prediction loop
def generate_caption(image, t=1, sample=False, max_len=20):
    """
    Generate caption for given image.
    if `sample` is True, we will sample next token from predicted probability distribution
    `t` is a temperature during that sampling,
    higher `t` causes more uniform-like distribution = more chaos.
    """
    # condition lstm on the image
    s.run(final_model.init_lstm,
          {final_model.input_images: [image]})

    # current caption
    # start with only START token
    caption = [vocab[START]]

    for _ in range(max_len):
        next_word_probs = s.run(final_model.one_step,
                                {final_model.current_word: [caption[-1]]})[0]
        next_word_probs = next_word_probs.ravel()

        # apply temperature
        next_word_probs = next_word_probs*(1/t) / np.sum(next_word_probs*(1/t))

        if sample:
            next_word = np.random.choice(range(len(vocab)), p=next_word_probs)
        else:
            next_word = np.argmax(next_word_probs)

        caption.append(next_word)
        if next_word == vocab[END]:
            break
```

```

        return list(map(vocab_inverse.get, caption))

In [42]: # look at validation prediction example
def apply_model_to_image_raw_bytes(raw):
    img = utils.decode_image_from_buf(raw)
    fig = plt.figure(figsize=(7, 7))
    plt.grid('off')
    plt.axis('off')
    plt.imshow(img)
    img = utils.crop_and_preprocess(img, (IMG_SIZE, IMG_SIZE), final_model.preprocess_f
    print(' '.join(generate_caption(img)[1:-1]))
    plt.show()

def show_valid_example(val_img_fns, example_idx=0):
    zf = zipfile.ZipFile("val2014_sample.zip")
    all_files = set(val_img_fns)
    found_files = list(filter(lambda x: x.filename.rsplit("/")[-1] in all_files, zf.fil
    example = found_files[example_idx]
    apply_model_to_image_raw_bytes(zf.read(example))

show_valid_example(val_img_fns, example_idx=100)

a baseball player holding a bat on a field

```



```
In [43]: # sample more images from validation
        for idx in np.random.choice(range(len(zipfile.ZipFile("val2014_sample.zip").filelist) -
            show_valid_example(val_img_fns, example_idx=idx)
            time.sleep(1)
```

a person on a ski slope with a ski lift



a woman holding a hot dog in front of a sandwich



a man is holding a cell phone in his hand



a large piece of cake with a large number of toppings



a bathroom with a sink and a sink



a bathroom with a sink and a sink



a man is throwing a frisbee in a park



a baseball player is getting ready to swing at a pitch



a piece of cake with a fork and a fork



a man is walking down a street with a surfboard



You can download any image from the Internet and apply your model to it!

```
In [49]: download_utils.download_file(  
        "https://h7f7z2r7.stackpathcdn.com/sites/default/files/images/articles/newyorkmain_  
        "newyorkmain_0.jpg"  
    )
```

HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443): Max retries exceeded with url:

Traceback (most recent call last):

```
File "/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py", line 595, in urlopen  
    self._prepare_proxy(conn)  
File "/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py", line 816, in _prepare  
    conn.connect()  
File "/opt/conda/lib/python3.6/site-packages/urllib3/connection.py", line 294, in connect  
    self._tunnel()  
File "/opt/conda/lib/python3.6/http/client.py", line 919, in _tunnel  
    message.strip()))
```

OSError: Tunnel connection failed: 403 Forbidden

During handling of the above exception, another exception occurred:

Traceback (most recent call last):

```
File "/opt/conda/lib/python3.6/site-packages/requests/adapters.py", line 440, in send  
    timeout=timeout  
File "/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py", line 639, in urlopen  
    _stacktrace=sys.exc_info()[2])  
File "/opt/conda/lib/python3.6/site-packages/urllib3/util/retry.py", line 388, in increment  
    raise MaxRetryError(_pool, url, error or ResponseError(cause))  
urllib3.exceptions.MaxRetryError: HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443
```

During handling of the above exception, another exception occurred:

Traceback (most recent call last):

```
File "../download_utils.py", line 21, in f_retry  
    return f(*args, **kwargs)  
File "../download_utils.py", line 39, in download_file  
    r = requests.get(url, stream=True)  
File "/opt/conda/lib/python3.6/site-packages/requests/api.py", line 72, in get  
    return request('get', url, params=params, **kwargs)  
File "/opt/conda/lib/python3.6/site-packages/requests/api.py", line 58, in request  
    return session.request(method=method, url=url, **kwargs)  
File "/opt/conda/lib/python3.6/site-packages/requests/sessions.py", line 508, in request  
    resp = self.send(prepare, **send_kwargs)  
File "/opt/conda/lib/python3.6/site-packages/requests/sessions.py", line 618, in send
```

```
    r = adapter.send(request, **kwargs)
File "/opt/conda/lib/python3.6/site-packages/requests/adapters.py", line 502, in send
    raise ProxyError(e, request=request)
requests.exceptions.ProxyError: HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443):
```

HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443): Max retries exceeded with url:

Traceback (most recent call last):

```
File "/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py", line 595, in urlopen
    self._prepare_proxy(conn)
File "/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py", line 816, in _prepare
    conn.connect()
File "/opt/conda/lib/python3.6/site-packages/urllib3/connection.py", line 294, in connect
    self._tunnel()
File "/opt/conda/lib/python3.6/http/client.py", line 919, in _tunnel
    message.strip()))
OSError: Tunnel connection failed: 403 Forbidden
```

During handling of the above exception, another exception occurred:

Traceback (most recent call last):

```
File "/opt/conda/lib/python3.6/site-packages/requests/adapters.py", line 440, in send
    timeout=timeout
File "/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py", line 639, in urlopen
    _stacktrace=sys.exc_info()[2])
File "/opt/conda/lib/python3.6/site-packages/urllib3/util/retry.py", line 388, in increment
    raise MaxRetryError(_pool, url, error or ResponseError(cause))
urllib3.exceptions.MaxRetryError: HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443)
```

During handling of the above exception, another exception occurred:

Traceback (most recent call last):

```
File "../download_utils.py", line 21, in f_retry
    return f(*args, **kwargs)
File "../download_utils.py", line 39, in download_file
    r = requests.get(url, stream=True)
File "/opt/conda/lib/python3.6/site-packages/requests/api.py", line 72, in get
    return request('get', url, params=params, **kwargs)
File "/opt/conda/lib/python3.6/site-packages/requests/api.py", line 58, in request
    return session.request(method=method, url=url, **kwargs)
File "/opt/conda/lib/python3.6/site-packages/requests/sessions.py", line 508, in request
    resp = self.send(prepare, **send_kwargs)
File "/opt/conda/lib/python3.6/site-packages/requests/sessions.py", line 618, in send
    r = adapter.send(request, **kwargs)
File "/opt/conda/lib/python3.6/site-packages/requests/adapters.py", line 502, in send
    raise ProxyError(e, request=request)
```

```
requests.exceptions.ProxyError: HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443):
```

```
HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443): Max retries exceeded with url:
```

```
Traceback (most recent call last):
```

```
File "/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py", line 595, in urlopen
    self._prepare_proxy(conn)
```

```
File "/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py", line 816, in _prepare
    conn.connect()
```

```
File "/opt/conda/lib/python3.6/site-packages/urllib3/connection.py", line 294, in connect
    self._tunnel()
```

```
File "/opt/conda/lib/python3.6/http/client.py", line 919, in _tunnel
    message.strip()))
```

```
OSError: Tunnel connection failed: 403 Forbidden
```

During handling of the above exception, another exception occurred:

```
Traceback (most recent call last):
```

```
File "/opt/conda/lib/python3.6/site-packages/requests/adapters.py", line 440, in send
    timeout=timeout
```

```
File "/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py", line 639, in urlopen
    _stacktrace=sys.exc_info()[2])
```

```
File "/opt/conda/lib/python3.6/site-packages/urllib3/util/retry.py", line 388, in increment
    raise MaxRetryError(_pool, url, error or ResponseError(cause))
```

```
urllib3.exceptions.MaxRetryError: HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443)
```

During handling of the above exception, another exception occurred:

```
Traceback (most recent call last):
```

```
File "../download_utils.py", line 21, in f_retry
    return f(*args, **kwargs)
```

```
File "../download_utils.py", line 39, in download_file
    r = requests.get(url, stream=True)
```

```
File "/opt/conda/lib/python3.6/site-packages/requests/api.py", line 72, in get
    return request('get', url, params=params, **kwargs)
```

```
File "/opt/conda/lib/python3.6/site-packages/requests/api.py", line 58, in request
    return session.request(method=method, url=url, **kwargs)
```

```
File "/opt/conda/lib/python3.6/site-packages/requests/sessions.py", line 508, in request
    resp = self.send(prepare, **send_kwargs)
```

```
File "/opt/conda/lib/python3.6/site-packages/requests/sessions.py", line 618, in send
    r = adapter.send(request, **kwargs)
```

```
File "/opt/conda/lib/python3.6/site-packages/requests/adapters.py", line 502, in send
    raise ProxyError(e, request=request)
```

```
requests.exceptions.ProxyError: HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443):
```

```

-----

OSError                                Traceback (most recent call last)

/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py in urlopen(self, method
594             if is_new_proxy_conn:
--> 595                 self._prepare_proxy(conn)
596

/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py in _prepare_proxy(self,
815
--> 816         conn.connect()
817

/opt/conda/lib/python3.6/site-packages/urllib3/connection.py in connect(self)
293             # self._tunnel_host below.
--> 294         self._tunnel()
295             # Mark this connection as not reusable

/opt/conda/lib/python3.6/http/client.py in _tunnel(self)
918             raise OSError("Tunnel connection failed: %d %s" % (code,
--> 919                                     message.strip()))
920         while True:

```

OSError: Tunnel connection failed: 403 Forbidden

During handling of the above exception, another exception occurred:

```

MaxRetryError                            Traceback (most recent call last)

/opt/conda/lib/python3.6/site-packages/requests/adapters.py in send(self, request, stream
439             retries=self.max_retries,
--> 440             timeout=timeout
441         )

/opt/conda/lib/python3.6/site-packages/urllib3/connectionpool.py in urlopen(self, method
638         retries = retries.increment(method, url, error=e, _pool=self,
--> 639                                     _stacktrace=sys.exc_info()[2])
640         retries.sleep()

```

```

/opt/conda/lib/python3.6/site-packages/urllib3/util/retry.py in increment(self, method,
387         if new_retry.is_exhausted():
--> 388             raise MaxRetryError(_pool, url, error or ResponseError(cause))
389

```

MaxRetryError: HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443): Max retr

During handling of the above exception, another exception occurred:

```

ProxyError                                Traceback (most recent call last)

<ipython-input-49-3d23f68a2091> in <module>()
    1 download_utils.download_file(
    2     "https://h7f7z2r7.stackpathcdn.com/sites/default/files/images/articles/newyorkma
----> 3     "newyorkmain_0.jpg"
    4 )

~/work/download_utils.py in f_retry(*args, **kwargs)
    28         mtries -= 1
    29         mdelay *= backoff
---> 30         return f(*args, **kwargs)
    31
    32         return f_retry # true decorator

~/work/download_utils.py in download_file(url, file_path)
    37 @retry(Exception)
    38 def download_file(url, file_path):
---> 39     r = requests.get(url, stream=True)
    40     total_size = int(r.headers.get('content-length'))
    41     bar = tqdm_utils.tqdm_notebook_failsafe(total=total_size, unit='B', unit_scale=T

/opt/conda/lib/python3.6/site-packages/requests/api.py in get(url, params, **kwargs)
    70
    71     kwargs.setdefault('allow_redirects', True)
---> 72     return request('get', url, params=params, **kwargs)
    73
    74

/opt/conda/lib/python3.6/site-packages/requests/api.py in request(method, url, **kwargs)
    56     # cases, and look like a memory leak in others.
    57     with sessions.Session() as session:

```

```

---> 58         return session.request(method=method, url=url, **kwargs)
      59
      60

/opt/conda/lib/python3.6/site-packages/requests/sessions.py in request(self, method, url,
506         }
507         send_kwargs.update(settings)
--> 508         resp = self.send(prepared_request, **send_kwargs)
      509
      510         return resp

/opt/conda/lib/python3.6/site-packages/requests/sessions.py in send(self, request, **kwargs)
616
617         # Send the request
--> 618         r = adapter.send(request, **kwargs)
      619
      620         # Total elapsed time of the request (approximately)

/opt/conda/lib/python3.6/site-packages/requests/adapters.py in send(self, request, stream, timeout, verify, cert, proxies)
500
501         if isinstance(e.reason, _ProxyError):
--> 502             raise ProxyError(e, request=request)
      503
      504         if isinstance(e.reason, _SSLError):

```

ProxyError: HTTPSConnectionPool(host='h7f7z2r7.stackpathcdn.com', port=443): Max retries

```
In [47]: apply_model_to_image_raw_bytes(open("portal-cake-10.jpg", "rb").read())
```

```
-----
FileNotFoundError                                Traceback (most recent call last)
```

```

<ipython-input-47-1e36111809c8> in <module>()
----> 1 apply_model_to_image_raw_bytes(open("http://www.bijouxandbits.com/wp-content/uploads

```

```
FileNotFoundError: [Errno 2] No such file or directory: 'http://www.bijouxandbits.com/wp
```

Now it's time to find 10 examples where your model works good and 10 examples where it fails!

You can use images from validation set as follows:


```
show_valid_example(val_img_fns, example_idx=...)
```

You can use images from the Internet as follows:

```
! wget ...  
apply_model_to_image_raw_bytes(open("...", "rb").read())
```

If you use these functions, the output will be embedded into your notebook and will be visible during peer review!

When you're done, download your notebook using "File" -> "Download as" -> "Notebook" and prepare that file for peer review!

```
In [ ]: ### YOUR EXAMPLES HERE ###
```

That's it!

Congratulations, you've trained your image captioning model and now can produce captions for any picture from the Internet!