

Coding Tutorial

November 20, 2020

```
In [1]: import tensorflow as tf
        print(tf.__version__)
```

2.0.0

1 Saving and loading models

Coding tutorials ##### Section ?? ##### Section ?? ##### Section ?? ##### Section ?? ##### Section ??

Saving and loading model weights

Load and inspect CIFAR-10 dataset The CIFAR-10 dataset consists of, in total, 60000 color images, each with one of 10 labels: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck. For an introduction and a download, see [this link](#).

```
In [2]: # Import the CIFAR-10 dataset and rescale the pixel values
```

```
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
x_train = x_train / 255.0
x_test = x_test / 255.0

# Use smaller subset -- speeds things up
x_train = x_train[:10000]
y_train = y_train[:10000]
x_test = x_test[:1000]
y_test = y_test[:1000]
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170500096/170498071 [=====] - 3s 0us/step

```
In [3]: # Plot the first 10 CIFAR-10 images
```

```
import matplotlib.pyplot as plt
```

```
fig, ax = plt.subplots(1, 10, figsize=(10, 1))
for i in range(10):
    ax[i].set_axis_off()
    ax[i].imshow(x_train[i])
```

Introduce two useful functions

In [3]: *# Introduce function to test model accuracy*

```
def get_test_accuracy(model, x_test, y_test):
    test_loss, test_acc = model.evaluate(x=x_test, y=y_test, verbose=0)
    print('accuracy: {acc:0.3f}'.format(acc=test_acc))
```

In [4]: *# Introduce function that creates a new instance of a simple CNN*

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D

def get_new_model():
    model = Sequential([
        Conv2D(filters=16, input_shape=(32, 32, 3), kernel_size=(3, 3),
            activation='relu', name='conv_1'),
        Conv2D(filters=8, kernel_size=(3, 3), activation='relu', name='conv_2'),
        MaxPooling2D(pool_size=(4, 4), name='pool_1'),
        Flatten(name='flatten'),
        Dense(units=32, activation='relu', name='dense_1'),
        Dense(units=10, activation='softmax', name='dense_2')
    ])
    model.compile(optimizer='adam',
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy'])
    return model
```

Create simple convolutional neural network classifier

In [6]: *# Create an instance of the model and show model summary*

```
model = get_new_model()
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv_1 (Conv2D)	(None, 30, 30, 16)	448
conv_2 (Conv2D)	(None, 28, 28, 8)	1160

pool_1 (MaxPooling2D)	(None, 7, 7, 8)	0

flatten (Flatten)	(None, 392)	0

dense_1 (Dense)	(None, 32)	12576

dense_2 (Dense)	(None, 10)	330
=====		
Total params: 14,514		
Trainable params: 14,514		
Non-trainable params: 0		

```
In [7]: # Test accuracy of the untrained model, around 10% (random)
        get_test_accuracy(model, x_test, y_test)
```

accuracy: 0.106

Train model with checkpoints

```
In [8]: from tensorflow.keras.callbacks import ModelCheckpoint
```

```
In [9]: # Create Tensorflow checkpoint object
        checkpoint_path = 'model_checkpoints/checkpoint'
        checkpoint = ModelCheckpoint(filepath = checkpoint_path, frequency = 'epoch',
                                     save_weights_only = True, verbose = 1)
```

```
In [10]: # Fit model, with simple checkpoint which saves (and overwrites) model weights every
         model.fit(x_train, y_train, epochs = 3, callbacks = [checkpoint])
```

Train on 10000 samples

Epoch 1/3

9984/10000 [=====>.] - ETA: 0s - loss: 1.9295 - accuracy: 0.3009

Epoch 00001: saving model to model_checkpoints/checkpoint

10000/10000 [=====] - 51s 5ms/sample - loss: 1.9291 - accuracy: 0.3009

Epoch 2/3

9984/10000 [=====>.] - ETA: 0s - loss: 1.5760 - accuracy: 0.4326

Epoch 00002: saving model to model_checkpoints/checkpoint

10000/10000 [=====] - 50s 5ms/sample - loss: 1.5757 - accuracy: 0.4326

Epoch 3/3

9984/10000 [=====>.] - ETA: 0s - loss: 1.4523 - accuracy: 0.4787

Epoch 00003: saving model to model_checkpoints/checkpoint

10000/10000 [=====] - 48s 5ms/sample - loss: 1.4530 - accuracy: 0.4787

```
Out[10]: <tensorflow.python.keras.callbacks.History at 0x7fa3f02e2240>
```

```
In [11]: # Have a look at what the checkpoint creates
!ls -lh model_checkpoints
```

```
total 184K
-rw-r--r-- 1 jovyan users 77 Nov 18 00:20 checkpoint
-rw-r--r-- 1 jovyan users 174K Nov 18 00:20 checkpoint.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Nov 18 00:20 checkpoint.index
```

```
In [14]: # Evaluate the performance of the trained model
get_test_accuracy(model, x_test, y_test)
```

```
accuracy: 0.482
```

Create new model, load weights

```
In [15]: # Create a new instance of the (initialised) model, accuracy around 10% again
model = get_new_model()
get_test_accuracy(model, x_test, y_test)
```

```
accuracy: 0.104
```

```
In [16]: # Load weights -- accuracy is the same as the trained model
model.load_weights(checkpoint_path)
get_test_accuracy(model, x_test, y_test)
```

```
accuracy: 0.482
```

Clear directory

```
In [17]: ! rm -r model_checkpoints
```

Model saving criteria

Create more customised checkpoint

```
In [5]: from tensorflow.keras.callbacks import ModelCheckpoint
```

```
In [6]: # Create Tensorflow checkpoint object with epoch and batch details
checkpoint_5000_path = 'model_checkpoints_5000/checkpoint_{epoch:02d}_{batch:04d}'
checkpoint_5000 = ModelCheckpoint(filepath = checkpoint_5000_path,
                                  save_weights_only = True,
                                  save_freq = 5000,
                                  verbose = 1)
```

```
In [7]: # Create and fit model with checkpoint
```

```
model = get_new_model()
model.fit(x_train, y_train, epochs = 3,
          validation_data = (x_test, y_test),
          batch_size = 10, callbacks = [checkpoint_5000])
```

Train on 10000 samples, validate on 1000 samples

Epoch 1/3

4990/10000 [=====>...] - ETA: 27s - loss: 2.0540 - accuracy: 0.2303

Epoch 00001: saving model to model_checkpoints_5000/checkpoint_01_0499

9990/10000 [=====>.] - ETA: 0s - loss: 1.8898 - accuracy: 0.2957

Epoch 00001: saving model to model_checkpoints_5000/checkpoint_01_0999

10000/10000 [=====] - 57s 6ms/sample - loss: 1.8896 - accuracy: 0.2957

Epoch 2/3

4990/10000 [=====>...] - ETA: 26s - loss: 1.5757 - accuracy: 0.4204

Epoch 00002: saving model to model_checkpoints_5000/checkpoint_02_0499

9980/10000 [=====>.] - ETA: 0s - loss: 1.5478 - accuracy: 0.4332

Epoch 00002: saving model to model_checkpoints_5000/checkpoint_02_0999

10000/10000 [=====] - 55s 5ms/sample - loss: 1.5478 - accuracy: 0.4332

Epoch 3/3

4990/10000 [=====>...] - ETA: 26s - loss: 1.4315 - accuracy: 0.4792

Epoch 00003: saving model to model_checkpoints_5000/checkpoint_03_0499

9990/10000 [=====>.] - ETA: 0s - loss: 1.4282 - accuracy: 0.4854

Epoch 00003: saving model to model_checkpoints_5000/checkpoint_03_0999

10000/10000 [=====] - 55s 5ms/sample - loss: 1.4281 - accuracy: 0.4854

```
Out[7]: <tensorflow.python.keras.callbacks.History at 0x7f5ac3b3b278>
```

```
In [8]: # Have a look at what the checkpoint creates
```

```
!ls -lh model_checkpoints_5000
```

total 1.1M

```
-rw-r--r-- 1 jovyan users 93 Nov 18 13:59 checkpoint
-rw-r--r-- 1 jovyan users 174K Nov 18 13:56 checkpoint_01_0499.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Nov 18 13:56 checkpoint_01_0499.index
-rw-r--r-- 1 jovyan users 174K Nov 18 13:57 checkpoint_01_0999.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Nov 18 13:57 checkpoint_01_0999.index
-rw-r--r-- 1 jovyan users 174K Nov 18 13:57 checkpoint_02_0499.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Nov 18 13:57 checkpoint_02_0499.index
-rw-r--r-- 1 jovyan users 174K Nov 18 13:58 checkpoint_02_0999.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Nov 18 13:58 checkpoint_02_0999.index
-rw-r--r-- 1 jovyan users 174K Nov 18 13:58 checkpoint_03_0499.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Nov 18 13:58 checkpoint_03_0499.index
-rw-r--r-- 1 jovyan users 174K Nov 18 13:59 checkpoint_03_0999.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Nov 18 13:59 checkpoint_03_0999.index
```

Work with model saving criteria

```
In [9]: # Use tiny training and test set -- will overfit!
```

```
x_train = x_train[:100]
y_train = y_train[:100]
x_test = x_test[:100]
y_test = y_test[:100]
```

```
In [10]: # Create a new instance of untrained model
```

```
model = get_new_model()
```

```
In [12]: # Create Tensorflow checkpoint object which monitors the validation accuracy
```

```
checkpoint_best_path = 'model_checkpoints_best/checkpoint'
checkpoint_best = ModelCheckpoint(filepath = checkpoint_best_path,
                                  save_weights_only = True,
                                  save_freq = 'epoch',
                                  monitor = 'val_accuracy',
                                  save_best_only = True,
                                  verbose = 1)
```

```
In [13]: # Fit the model and save only the weights with the highest validation accuracy
```

```
history = model.fit(x_train, y_train, epochs = 50,
                    validation_data = (x_test, y_test),
                    batch_size = 10,
                    callbacks = [checkpoint_best],
                    verbose = 0)
```

```
Epoch 00001: val_accuracy improved from -inf to 0.12000, saving model to model_checkpoints_best
```

```
Epoch 00002: val_accuracy did not improve from 0.12000
```

```
Epoch 00003: val_accuracy did not improve from 0.12000
```

```
Epoch 00004: val_accuracy did not improve from 0.12000
```

```
Epoch 00005: val_accuracy improved from 0.12000 to 0.14000, saving model to model_checkpoints_l
```

```
Epoch 00006: val_accuracy did not improve from 0.14000
```

```
Epoch 00007: val_accuracy did not improve from 0.14000
```

```
Epoch 00008: val_accuracy did not improve from 0.14000
```

```
Epoch 00009: val_accuracy did not improve from 0.14000
```

```
Epoch 00010: val_accuracy improved from 0.14000 to 0.15000, saving model to model_checkpoints_l
```

Epoch 00011: val_accuracy did not improve from 0.15000

Epoch 00012: val_accuracy improved from 0.15000 to 0.17000, saving model to model_checkpoints_1

Epoch 00013: val_accuracy did not improve from 0.17000

Epoch 00014: val_accuracy did not improve from 0.17000

Epoch 00015: val_accuracy did not improve from 0.17000

Epoch 00016: val_accuracy did not improve from 0.17000

Epoch 00017: val_accuracy did not improve from 0.17000

Epoch 00018: val_accuracy improved from 0.17000 to 0.20000, saving model to model_checkpoints_1

Epoch 00019: val_accuracy did not improve from 0.20000

Epoch 00020: val_accuracy did not improve from 0.20000

Epoch 00021: val_accuracy did not improve from 0.20000

Epoch 00022: val_accuracy did not improve from 0.20000

Epoch 00023: val_accuracy did not improve from 0.20000

Epoch 00024: val_accuracy did not improve from 0.20000

Epoch 00025: val_accuracy did not improve from 0.20000

Epoch 00026: val_accuracy did not improve from 0.20000

Epoch 00027: val_accuracy did not improve from 0.20000

Epoch 00028: val_accuracy did not improve from 0.20000

Epoch 00029: val_accuracy did not improve from 0.20000

Epoch 00030: val_accuracy did not improve from 0.20000

Epoch 00031: val_accuracy did not improve from 0.20000

Epoch 00032: val_accuracy did not improve from 0.20000

Epoch 00033: val_accuracy did not improve from 0.20000

Epoch 00034: val_accuracy did not improve from 0.20000


```
Epoch 00035: val_accuracy did not improve from 0.20000
Epoch 00036: val_accuracy did not improve from 0.20000
Epoch 00037: val_accuracy improved from 0.20000 to 0.21000, saving model to model_checkpoints_1
Epoch 00038: val_accuracy did not improve from 0.21000
Epoch 00039: val_accuracy did not improve from 0.21000
Epoch 00040: val_accuracy improved from 0.21000 to 0.22000, saving model to model_checkpoints_1
Epoch 00041: val_accuracy improved from 0.22000 to 0.23000, saving model to model_checkpoints_1
Epoch 00042: val_accuracy did not improve from 0.23000
Epoch 00043: val_accuracy did not improve from 0.23000
Epoch 00044: val_accuracy did not improve from 0.23000
Epoch 00045: val_accuracy did not improve from 0.23000
Epoch 00046: val_accuracy did not improve from 0.23000
Epoch 00047: val_accuracy did not improve from 0.23000
Epoch 00048: val_accuracy did not improve from 0.23000
Epoch 00049: val_accuracy did not improve from 0.23000
Epoch 00050: val_accuracy did not improve from 0.23000
```

```
In [14]: # Plot training and testing curves
```

```
import pandas as pd

df = pd.DataFrame(history.history)
df.plot(y=['accuracy', 'val_accuracy'])
```

```
Out[14]: <matplotlib.axes._subplots.AxesSubplot at 0x7f5b69d113c8>
```

```
In [15]: # Inspect the checkpoint directory
         !ls -lh model_checkpoints_best
```

```
total 184K
-rw-r--r-- 1 jovyan users 77 Nov 18 14:22 checkpoint
-rw-r--r-- 1 jovyan users 174K Nov 18 14:22 checkpoint.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.0K Nov 18 14:22 checkpoint.index
```

```
In [18]: # Create a new model with the saved weights
         new_model = get_new_model()
         new_model.load_weights(checkpoint_best_path)
         get_test_accuracy(new_model, x_test, y_test)
```

```
accuracy: 0.230
```

Clear directory

```
In [19]: ! rm -r model_checkpoints_5000 model_checkpoints_best
```

Saving the entire model

Create checkpoint that saves whole model, not just weights

```
In [5]: from tensorflow.keras.callbacks import ModelCheckpoint
```

```
In [6]: # Create Tensorflow checkpoint object
checkpoint_path = 'model_checkpoints'
checkpoint = ModelCheckpoint(filepath = checkpoint_path,
                             save_weights_only = False,
                             frequency = 'epoch',
                             verbose = 1)
```

```
In [7]: # Create and fit model with checkpoint
model = get_new_model()
model.fit(x_train, y_train, epochs = 3, callbacks = [checkpoint])
```

Train on 10000 samples

Epoch 1/3

9984/10000 [=====>.] - ETA: 0s - loss: 2.0021 - accuracy: 0.2648

Epoch 00001: saving model to model_checkpoints

WARNING:tensorflow:From /opt/conda/lib/python3.7/site-packages/tensorflow_core/python/ops/resource_ops.py:165: tf.nn.conv2d is deprecated and will be removed in a future version.

Instructions for updating:

If using Keras pass *_constraint arguments to layers.

INFO:tensorflow:Assets written to: model_checkpoints/assets

10000/10000 [=====] - 47s 5ms/sample - loss: 2.0015 - accuracy: 0.2650

Epoch 2/3

9984/10000 [=====>.] - ETA: 0s - loss: 1.6207 - accuracy: 0.4116

Epoch 00002: saving model to model_checkpoints

INFO:tensorflow:Assets written to: model_checkpoints/assets

10000/10000 [=====] - 45s 5ms/sample - loss: 1.6206 - accuracy: 0.4116

Epoch 3/3

9984/10000 [=====>.] - ETA: 0s - loss: 1.4863 - accuracy: 0.4649

Epoch 00003: saving model to model_checkpoints

INFO:tensorflow:Assets written to: model_checkpoints/assets

10000/10000 [=====] - 44s 4ms/sample - loss: 1.4865 - accuracy: 0.4649

```
Out[7]: <tensorflow.python.keras.callbacks.History at 0x7f3ebc403f98>
```

Inspect what the checkpoint has created

```
In [8]: # Have a look at what the checkpoint creates
!ls -lh model_checkpoints
```

total 128K

drwxr-xr-x 2 jovyan users 6.0K Nov 19 03:40 assets

-rw-r--r-- 1 jovyan users 118K Nov 19 03:42 saved_model.pb

drwxr-xr-x 2 jovyan users 6.0K Nov 19 03:42 variables

```
In [9]: # Enter variables directory
        !ls -lh model_checkpoints/variables

total 184K
-rw-r--r-- 1 jovyan users 177K Nov 19 03:42 variables.data-00000-of-00001
-rw-r--r-- 1 jovyan users 2.1K Nov 19 03:42 variables.index

In [10]: # Get the model's test accuracy
         get_test_accuracy(model, x_test, y_test)

accuracy: 0.470
```

Create new model from scratch

```
In [11]: # Delete model
         del model

In [13]: from tensorflow.keras.models import load_model

In [15]: # Reload model from scratch
         model = load_model(checkpoint_path)
         get_test_accuracy(model, x_test, y_test)

accuracy: 0.470
```

Use the .h5 format to save model

```
In [16]: # Save the model in .h5 format
         model.save('my_model.h5')

In [17]: # Inspect .h5 file
         !ls -lh my_model.h5

-rw-r--r-- 1 jovyan users 77K Nov 19 03:47 my_model.h5

In [18]: # Delete model
         del model

In [20]: # Reload model from scratch
         model = load_model('my_model.h5')
         get_test_accuracy(model, x_test, y_test)

accuracy: 0.470
```

Clear directory

```
In [21]: ! rm -r model_checkpoints
        ! rm my_model.h5
```

Loading pre-trained Keras models

Import and build Keras ResNet50 model Today we'll be using the ResNet50 model designed by a team at Microsoft Research, available through Keras applications. Please see the description on the [Keras applications page](#) for details. If you continue using it, please cite it properly! The paper it comes from is:

Kaiming He, Xiangyu Zhang, Shaoqing Ren, Jian Sun. "Deep Residual Learning for Image Recognition", 2015.

This model takes a long time to download on the Coursera platform, so it is pre-downloaded in your workspace and saved in Keras HDF5 format. If you want to import it on your personal machine, use the following code:

```
from tensorflow.keras.applications import ResNet50
model = ResNet50(weights='imagenet')
```

In this coding tutorial, you will instead load the model directly from disk.

```
In [2]: from tensorflow.keras.models import load_model
```

```
In [3]: # Build Keras ResNet50 model
        model = load_model('models/Keras_ResNet50.h5')
        model.summary()
```

```
WARNING:tensorflow:No training configuration found in save file: the model was *not* compiled.
Model: "resnet50"
```

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[(None, 224, 224, 3)]	0	
conv1_pad (ZeroPadding2D)	(None, 230, 230, 3)	0	input_1[0][0]
conv1_conv (Conv2D)	(None, 112, 112, 64)	9472	conv1_pad[0][0]
conv1_bn (BatchNormalization)	(None, 112, 112, 64)	256	conv1_conv[0][0]
conv1_relu (Activation)	(None, 112, 112, 64)	0	conv1_bn[0][0]
pool1_pad (ZeroPadding2D)	(None, 114, 114, 64)	0	conv1_relu[0][0]
pool1_pool (MaxPooling2D)	(None, 56, 56, 64)	0	pool1_pad[0][0]

conv2_block1_1_conv (Conv2D)	(None, 56, 56, 64)	4160	pool1_pool[0][0]
conv2_block1_1_bn (BatchNormali	(None, 56, 56, 64)	256	conv2_block1_1_conv[0][0]
conv2_block1_1_relu (Activation	(None, 56, 56, 64)	0	conv2_block1_1_bn[0][0]
conv2_block1_2_conv (Conv2D)	(None, 56, 56, 64)	36928	conv2_block1_1_relu[0][0]
conv2_block1_2_bn (BatchNormali	(None, 56, 56, 64)	256	conv2_block1_2_conv[0][0]
conv2_block1_2_relu (Activation	(None, 56, 56, 64)	0	conv2_block1_2_bn[0][0]
conv2_block1_0_conv (Conv2D)	(None, 56, 56, 256)	16640	pool1_pool[0][0]
conv2_block1_3_conv (Conv2D)	(None, 56, 56, 256)	16640	conv2_block1_2_relu[0][0]
conv2_block1_0_bn (BatchNormali	(None, 56, 56, 256)	1024	conv2_block1_0_conv[0][0]
conv2_block1_3_bn (BatchNormali	(None, 56, 56, 256)	1024	conv2_block1_3_conv[0][0]
conv2_block1_add (Add)	(None, 56, 56, 256)	0	conv2_block1_0_bn[0][0] conv2_block1_3_bn[0][0]
conv2_block1_out (Activation)	(None, 56, 56, 256)	0	conv2_block1_add[0][0]
conv2_block2_1_conv (Conv2D)	(None, 56, 56, 64)	16448	conv2_block1_out[0][0]
conv2_block2_1_bn (BatchNormali	(None, 56, 56, 64)	256	conv2_block2_1_conv[0][0]
conv2_block2_1_relu (Activation	(None, 56, 56, 64)	0	conv2_block2_1_bn[0][0]
conv2_block2_2_conv (Conv2D)	(None, 56, 56, 64)	36928	conv2_block2_1_relu[0][0]
conv2_block2_2_bn (BatchNormali	(None, 56, 56, 64)	256	conv2_block2_2_conv[0][0]
conv2_block2_2_relu (Activation	(None, 56, 56, 64)	0	conv2_block2_2_bn[0][0]
conv2_block2_3_conv (Conv2D)	(None, 56, 56, 256)	16640	conv2_block2_2_relu[0][0]
conv2_block2_3_bn (BatchNormali	(None, 56, 56, 256)	1024	conv2_block2_3_conv[0][0]
conv2_block2_add (Add)	(None, 56, 56, 256)	0	conv2_block1_out[0][0] conv2_block2_3_bn[0][0]
conv2_block2_out (Activation)	(None, 56, 56, 256)	0	conv2_block2_add[0][0]
conv2_block3_1_conv (Conv2D)	(None, 56, 56, 64)	16448	conv2_block2_out[0][0]

conv2_block3_1_bn (BatchNormali	(None, 56, 56, 64)	256	conv2_block3_1_conv[0][0]
conv2_block3_1_relu (Activation	(None, 56, 56, 64)	0	conv2_block3_1_bn[0][0]
conv2_block3_2_conv (Conv2D)	(None, 56, 56, 64)	36928	conv2_block3_1_relu[0][0]
conv2_block3_2_bn (BatchNormali	(None, 56, 56, 64)	256	conv2_block3_2_conv[0][0]
conv2_block3_2_relu (Activation	(None, 56, 56, 64)	0	conv2_block3_2_bn[0][0]
conv2_block3_3_conv (Conv2D)	(None, 56, 56, 256)	16640	conv2_block3_2_relu[0][0]
conv2_block3_3_bn (BatchNormali	(None, 56, 56, 256)	1024	conv2_block3_3_conv[0][0]
conv2_block3_add (Add)	(None, 56, 56, 256)	0	conv2_block2_out[0][0] conv2_block3_3_bn[0][0]
conv2_block3_out (Activation)	(None, 56, 56, 256)	0	conv2_block3_add[0][0]
conv3_block1_1_conv (Conv2D)	(None, 28, 28, 128)	32896	conv2_block3_out[0][0]
conv3_block1_1_bn (BatchNormali	(None, 28, 28, 128)	512	conv3_block1_1_conv[0][0]
conv3_block1_1_relu (Activation	(None, 28, 28, 128)	0	conv3_block1_1_bn[0][0]
conv3_block1_2_conv (Conv2D)	(None, 28, 28, 128)	147584	conv3_block1_1_relu[0][0]
conv3_block1_2_bn (BatchNormali	(None, 28, 28, 128)	512	conv3_block1_2_conv[0][0]
conv3_block1_2_relu (Activation	(None, 28, 28, 128)	0	conv3_block1_2_bn[0][0]
conv3_block1_0_conv (Conv2D)	(None, 28, 28, 512)	131584	conv2_block3_out[0][0]
conv3_block1_3_conv (Conv2D)	(None, 28, 28, 512)	66048	conv3_block1_2_relu[0][0]
conv3_block1_0_bn (BatchNormali	(None, 28, 28, 512)	2048	conv3_block1_0_conv[0][0]
conv3_block1_3_bn (BatchNormali	(None, 28, 28, 512)	2048	conv3_block1_3_conv[0][0]
conv3_block1_add (Add)	(None, 28, 28, 512)	0	conv3_block1_0_bn[0][0] conv3_block1_3_bn[0][0]
conv3_block1_out (Activation)	(None, 28, 28, 512)	0	conv3_block1_add[0][0]
conv3_block2_1_conv (Conv2D)	(None, 28, 28, 128)	65664	conv3_block1_out[0][0]
conv3_block2_1_bn (BatchNormali	(None, 28, 28, 128)	512	conv3_block2_1_conv[0][0]

conv3_block2_1_relu	(Activation	(None, 28, 28, 128)	0	conv3_block2_1_bn[0][0]
conv3_block2_2_conv	(Conv2D)	(None, 28, 28, 128)	147584	conv3_block2_1_relu[0][0]
conv3_block2_2_bn	(BatchNormali	(None, 28, 28, 128)	512	conv3_block2_2_conv[0][0]
conv3_block2_2_relu	(Activation	(None, 28, 28, 128)	0	conv3_block2_2_bn[0][0]
conv3_block2_3_conv	(Conv2D)	(None, 28, 28, 512)	66048	conv3_block2_2_relu[0][0]
conv3_block2_3_bn	(BatchNormali	(None, 28, 28, 512)	2048	conv3_block2_3_conv[0][0]
conv3_block2_add	(Add)	(None, 28, 28, 512)	0	conv3_block1_out[0][0] conv3_block2_3_bn[0][0]
conv3_block2_out	(Activation)	(None, 28, 28, 512)	0	conv3_block2_add[0][0]
conv3_block3_1_conv	(Conv2D)	(None, 28, 28, 128)	65664	conv3_block2_out[0][0]
conv3_block3_1_bn	(BatchNormali	(None, 28, 28, 128)	512	conv3_block3_1_conv[0][0]
conv3_block3_1_relu	(Activation	(None, 28, 28, 128)	0	conv3_block3_1_bn[0][0]
conv3_block3_2_conv	(Conv2D)	(None, 28, 28, 128)	147584	conv3_block3_1_relu[0][0]
conv3_block3_2_bn	(BatchNormali	(None, 28, 28, 128)	512	conv3_block3_2_conv[0][0]
conv3_block3_2_relu	(Activation	(None, 28, 28, 128)	0	conv3_block3_2_bn[0][0]
conv3_block3_3_conv	(Conv2D)	(None, 28, 28, 512)	66048	conv3_block3_2_relu[0][0]
conv3_block3_3_bn	(BatchNormali	(None, 28, 28, 512)	2048	conv3_block3_3_conv[0][0]
conv3_block3_add	(Add)	(None, 28, 28, 512)	0	conv3_block2_out[0][0] conv3_block3_3_bn[0][0]
conv3_block3_out	(Activation)	(None, 28, 28, 512)	0	conv3_block3_add[0][0]
conv3_block4_1_conv	(Conv2D)	(None, 28, 28, 128)	65664	conv3_block3_out[0][0]
conv3_block4_1_bn	(BatchNormali	(None, 28, 28, 128)	512	conv3_block4_1_conv[0][0]
conv3_block4_1_relu	(Activation	(None, 28, 28, 128)	0	conv3_block4_1_bn[0][0]
conv3_block4_2_conv	(Conv2D)	(None, 28, 28, 128)	147584	conv3_block4_1_relu[0][0]
conv3_block4_2_bn	(BatchNormali	(None, 28, 28, 128)	512	conv3_block4_2_conv[0][0]

conv3_block4_2_relu (Activation	(None, 28, 28, 128)	0	conv3_block4_2_bn[0][0]
conv3_block4_3_conv (Conv2D)	(None, 28, 28, 512)	66048	conv3_block4_2_relu[0][0]
conv3_block4_3_bn (BatchNormali	(None, 28, 28, 512)	2048	conv3_block4_3_conv[0][0]
conv3_block4_add (Add)	(None, 28, 28, 512)	0	conv3_block3_out[0][0] conv3_block4_3_bn[0][0]
conv3_block4_out (Activation)	(None, 28, 28, 512)	0	conv3_block4_add[0][0]
conv4_block1_1_conv (Conv2D)	(None, 14, 14, 256)	131328	conv3_block4_out[0][0]
conv4_block1_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block1_1_conv[0][0]
conv4_block1_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block1_1_bn[0][0]
conv4_block1_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block1_1_relu[0][0]
conv4_block1_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block1_2_conv[0][0]
conv4_block1_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block1_2_bn[0][0]
conv4_block1_0_conv (Conv2D)	(None, 14, 14, 1024)	525312	conv3_block4_out[0][0]
conv4_block1_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block1_2_relu[0][0]
conv4_block1_0_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block1_0_conv[0][0]
conv4_block1_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block1_3_conv[0][0]
conv4_block1_add (Add)	(None, 14, 14, 1024)	0	conv4_block1_0_bn[0][0] conv4_block1_3_bn[0][0]
conv4_block1_out (Activation)	(None, 14, 14, 1024)	0	conv4_block1_add[0][0]
conv4_block2_1_conv (Conv2D)	(None, 14, 14, 256)	262400	conv4_block1_out[0][0]
conv4_block2_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block2_1_conv[0][0]
conv4_block2_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block2_1_bn[0][0]
conv4_block2_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block2_1_relu[0][0]
conv4_block2_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block2_2_conv[0][0]
conv4_block2_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block2_2_bn[0][0]

conv4_block2_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block2_2_relu[0][0]
conv4_block2_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block2_3_conv[0][0]
conv4_block2_add (Add)	(None, 14, 14, 1024)	0	conv4_block1_out[0][0] conv4_block2_3_bn[0][0]
conv4_block2_out (Activation)	(None, 14, 14, 1024)	0	conv4_block2_add[0][0]
conv4_block3_1_conv (Conv2D)	(None, 14, 14, 256)	262400	conv4_block2_out[0][0]
conv4_block3_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block3_1_conv[0][0]
conv4_block3_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block3_1_bn[0][0]
conv4_block3_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block3_1_relu[0][0]
conv4_block3_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block3_2_conv[0][0]
conv4_block3_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block3_2_bn[0][0]
conv4_block3_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block3_2_relu[0][0]
conv4_block3_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block3_3_conv[0][0]
conv4_block3_add (Add)	(None, 14, 14, 1024)	0	conv4_block2_out[0][0] conv4_block3_3_bn[0][0]
conv4_block3_out (Activation)	(None, 14, 14, 1024)	0	conv4_block3_add[0][0]
conv4_block4_1_conv (Conv2D)	(None, 14, 14, 256)	262400	conv4_block3_out[0][0]
conv4_block4_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block4_1_conv[0][0]
conv4_block4_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block4_1_bn[0][0]
conv4_block4_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block4_1_relu[0][0]
conv4_block4_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block4_2_conv[0][0]
conv4_block4_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block4_2_bn[0][0]
conv4_block4_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block4_2_relu[0][0]
conv4_block4_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block4_3_conv[0][0]
conv4_block4_add (Add)	(None, 14, 14, 1024)	0	conv4_block3_out[0][0] conv4_block4_3_bn[0][0]

conv4_block4_out (Activation)	(None, 14, 14, 1024)	0	conv4_block4_add[0][0]
conv4_block5_1_conv (Conv2D)	(None, 14, 14, 256)	262400	conv4_block4_out[0][0]
conv4_block5_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block5_1_conv[0][0]
conv4_block5_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block5_1_bn[0][0]
conv4_block5_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block5_1_relu[0][0]
conv4_block5_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block5_2_conv[0][0]
conv4_block5_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block5_2_bn[0][0]
conv4_block5_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block5_2_relu[0][0]
conv4_block5_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block5_3_conv[0][0]
conv4_block5_add (Add)	(None, 14, 14, 1024)	0	conv4_block4_out[0][0] conv4_block5_3_bn[0][0]
conv4_block5_out (Activation)	(None, 14, 14, 1024)	0	conv4_block5_add[0][0]
conv4_block6_1_conv (Conv2D)	(None, 14, 14, 256)	262400	conv4_block5_out[0][0]
conv4_block6_1_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block6_1_conv[0][0]
conv4_block6_1_relu (Activation	(None, 14, 14, 256)	0	conv4_block6_1_bn[0][0]
conv4_block6_2_conv (Conv2D)	(None, 14, 14, 256)	590080	conv4_block6_1_relu[0][0]
conv4_block6_2_bn (BatchNormali	(None, 14, 14, 256)	1024	conv4_block6_2_conv[0][0]
conv4_block6_2_relu (Activation	(None, 14, 14, 256)	0	conv4_block6_2_bn[0][0]
conv4_block6_3_conv (Conv2D)	(None, 14, 14, 1024)	263168	conv4_block6_2_relu[0][0]
conv4_block6_3_bn (BatchNormali	(None, 14, 14, 1024)	4096	conv4_block6_3_conv[0][0]
conv4_block6_add (Add)	(None, 14, 14, 1024)	0	conv4_block5_out[0][0] conv4_block6_3_bn[0][0]
conv4_block6_out (Activation)	(None, 14, 14, 1024)	0	conv4_block6_add[0][0]
conv5_block1_1_conv (Conv2D)	(None, 7, 7, 512)	524800	conv4_block6_out[0][0]
conv5_block1_1_bn (BatchNormali	(None, 7, 7, 512)	2048	conv5_block1_1_conv[0][0]

conv5_block1_1_relu	(Activation	(None, 7, 7, 512)	0	conv5_block1_1_bn[0][0]
conv5_block1_2_conv	(Conv2D)	(None, 7, 7, 512)	2359808	conv5_block1_1_relu[0][0]
conv5_block1_2_bn	(BatchNormali	(None, 7, 7, 512)	2048	conv5_block1_2_conv[0][0]
conv5_block1_2_relu	(Activation	(None, 7, 7, 512)	0	conv5_block1_2_bn[0][0]
conv5_block1_0_conv	(Conv2D)	(None, 7, 7, 2048)	2099200	conv4_block6_out[0][0]
conv5_block1_3_conv	(Conv2D)	(None, 7, 7, 2048)	1050624	conv5_block1_2_relu[0][0]
conv5_block1_0_bn	(BatchNormali	(None, 7, 7, 2048)	8192	conv5_block1_0_conv[0][0]
conv5_block1_3_bn	(BatchNormali	(None, 7, 7, 2048)	8192	conv5_block1_3_conv[0][0]
conv5_block1_add	(Add)	(None, 7, 7, 2048)	0	conv5_block1_0_bn[0][0] conv5_block1_3_bn[0][0]
conv5_block1_out	(Activation)	(None, 7, 7, 2048)	0	conv5_block1_add[0][0]
conv5_block2_1_conv	(Conv2D)	(None, 7, 7, 512)	1049088	conv5_block1_out[0][0]
conv5_block2_1_bn	(BatchNormali	(None, 7, 7, 512)	2048	conv5_block2_1_conv[0][0]
conv5_block2_1_relu	(Activation	(None, 7, 7, 512)	0	conv5_block2_1_bn[0][0]
conv5_block2_2_conv	(Conv2D)	(None, 7, 7, 512)	2359808	conv5_block2_1_relu[0][0]
conv5_block2_2_bn	(BatchNormali	(None, 7, 7, 512)	2048	conv5_block2_2_conv[0][0]
conv5_block2_2_relu	(Activation	(None, 7, 7, 512)	0	conv5_block2_2_bn[0][0]
conv5_block2_3_conv	(Conv2D)	(None, 7, 7, 2048)	1050624	conv5_block2_2_relu[0][0]
conv5_block2_3_bn	(BatchNormali	(None, 7, 7, 2048)	8192	conv5_block2_3_conv[0][0]
conv5_block2_add	(Add)	(None, 7, 7, 2048)	0	conv5_block1_out[0][0] conv5_block2_3_bn[0][0]
conv5_block2_out	(Activation)	(None, 7, 7, 2048)	0	conv5_block2_add[0][0]
conv5_block3_1_conv	(Conv2D)	(None, 7, 7, 512)	1049088	conv5_block2_out[0][0]
conv5_block3_1_bn	(BatchNormali	(None, 7, 7, 512)	2048	conv5_block3_1_conv[0][0]
conv5_block3_1_relu	(Activation	(None, 7, 7, 512)	0	conv5_block3_1_bn[0][0]

conv5_block3_2_conv (Conv2D)	(None, 7, 7, 512)	2359808	conv5_block3_1_relu[0][0]
conv5_block3_2_bn (BatchNormali	(None, 7, 7, 512)	2048	conv5_block3_2_conv[0][0]
conv5_block3_2_relu (Activation	(None, 7, 7, 512)	0	conv5_block3_2_bn[0][0]
conv5_block3_3_conv (Conv2D)	(None, 7, 7, 2048)	1050624	conv5_block3_2_relu[0][0]
conv5_block3_3_bn (BatchNormali	(None, 7, 7, 2048)	8192	conv5_block3_3_conv[0][0]
conv5_block3_add (Add)	(None, 7, 7, 2048)	0	conv5_block2_out[0][0] conv5_block3_3_bn[0][0]
conv5_block3_out (Activation)	(None, 7, 7, 2048)	0	conv5_block3_add[0][0]
avg_pool (GlobalAveragePooling2	(None, 2048)	0	conv5_block3_out[0][0]
probs (Dense)	(None, 1000)	2049000	avg_pool[0][0]

=====
 Total params: 25,636,712
 Trainable params: 25,583,592
 Non-trainable params: 53,120
 =====

Import and preprocess 3 sample images

In [4]: *# Import 3 sample ImageNet images*

```

from tensorflow.keras.preprocessing.image import load_img

lemon_img = load_img('data/lemon.jpg', target_size=(224, 224))
viaduct_img = load_img('data/viaduct.jpg', target_size=(224, 224))
water_tower_img = load_img('data/water_tower.jpg', target_size=(224, 224))

```

Use ResNet50 model to classify images

In [5]: *# Useful function: presents top 5 predictions and probabilities*

```

from tensorflow.keras.preprocessing.image import img_to_array
from tensorflow.keras.applications.resnet50 import preprocess_input, decode_predictions
import numpy as np
import pandas as pd

def get_top_5_predictions(img):
    x = img_to_array(img)[np.newaxis, ...]
    x = preprocess_input(x)

```

```

preds = decode_predictions(model.predict(x), top=5)
top_preds = pd.DataFrame(columns=['prediction', 'probability'],
                           index=np.arange(5)+1)

for i in range(5):
    top_preds.loc[i+1, 'prediction'] = preds[0][i][1]
    top_preds.loc[i+1, 'probability'] = preds[0][i][2]
return top_preds

```

Image 1: lemon

In [6]: # Display image
lemon_img

Out [6]:



In [8]: # Display top 5 predictions
get_top_5_predictions(lemon_img)

Downloading data from https://storage.googleapis.com/download.tensorflow.org/data/imagenet_classification_2015_synset_words.txt
40960/35363 [=====] - 0s 0us/step

Out [8]:

	prediction	probability
1	consomme	0.227801
2	lemon	0.221758
3	egg nog	0.151335
4	ladle	0.0400883
5	spotlight	0.0291972

Image 2: viaduct

```
In [9]: # Display image  
viaduct_img
```

Out [9]:



```
In [10]: # Display top 5 predictions  
get_top_5_predictions(viaduct_img)
```

Out [10]:

	prediction	probability
1	vault	0.367951
2	prison	0.111742
3	viaduct	0.110738
4	window_screen	0.0882307
5	fire_screen	0.0206752

Image 3: water tower

```
In [11]: # Display image  
water_tower_img
```

Out [11]:



```
In [12]: # Display top 5 predictions
         get_top_5_predictions(water_tower_img)
```

```
Out[12]:
```

	prediction	probability
1	solar_dish	0.384681
2	ladle	0.196179
3	planetarium	0.116682
4	strainer	0.04999
5	jigsaw_puzzle	0.0219473

Tensorflow Hub modules

Import and build Tensorflow Hub MobileNet v1 model Today we'll be using Google's MobileNet v1 model, available on Tensorflow Hub. Please see the description on the [Tensorflow Hub page](#) for details on it's architecture, how it's trained, and the reference. If you continue using it, please cite it properly! The paper it comes from is:

Andrew G. Howard, Menglong Zhu, Bo Chen, Dmitry Kalenichenko, Weijun Wang, Tobias Weyand, Marco Andreetto, Hartwig Adam: "MobileNets: Efficient Convolutional Neural Networks for Mobile Vision Applications", 2017.

This model takes a long time to download on the Coursera platform, so it is pre-downloaded in your workspace and saved in Tensorflow SavedModel format. If you want to import it on your personal machine, use the following code:

```
module_url = "https://tfhub.dev/google/imagenet/mobilenet_v1_050_160/classification/4"
model = Sequential([hub.KerasLayer(module_url)])
model.build(input_shape=[None, 160, 160, 3])
```

In this coding tutorial, you will instead load the model directly from disk.

```
In [4]: import tensorflow_hub as hub
        from tensorflow.keras.models import load_model, Sequential
```

```
In [17]: import numpy as np
         import pandas as pd
```

```
In [5]: # Build Google's Mobilenet v1 model
        module = load_model('models/Tensorflow_MobileNet_v1')
        model = Sequential(hub.KerasLayer(module))
        model.build(input_shape = [None, 160, 160, 3])
        model.summary()
```

Model: "sequential"

```
-----
Layer (type)                 Output Shape          Param #
=====
keras_layer (KerasLayer)     multiple              1343049
=====
Total params: 1,343,049
Trainable params: 0
Non-trainable params: 1,343,049
-----
```

Use MobileNet model to classify images

```
In [11]: # Import and preprocess 3 sample ImageNet images
```

```
        from tensorflow.keras.preprocessing.image import load_img, img_to_array

        lemon_img = load_img("data/lemon.jpg", target_size=(160, 160))
        viaduct_img = load_img("data/viaduct.jpg", target_size=(160, 160))
        water_tower_img = load_img("data/water_tower.jpg", target_size=(160, 160))
```

```
In [13]: # Read in categories text file
```

```
        with open('data/imagenet_categories.txt') as txt_file:
            categories = txt_file.read().splitlines()
```

```
In [14]: # Useful function: presents top 5 predictions
```

```
        import pandas as pd

        def get_top_5_predictions(img):
            x = img_to_array(img)[np.newaxis, ...] / 255.0
            preds = model.predict(x)
            top_preds = pd.DataFrame(columns=['prediction'],
```



```

                                index=np.arange(5)+1)
sorted_index = np.argsort(-preds[0])
for i in range(5):
    ith_pred = categories[sorted_index[i]]
    top_preds.loc[i+1, 'prediction'] = ith_pred

return top_preds

```

Image 1: lemon

In [15]: lemon_img

Out[15]:



In [18]: get_top_5_predictions(lemon_img)

```

Out[18]:
      prediction
1      shower cap
2           tray
3         candle
4      brassiere
5 African chameleon

```

Image 2: viaduct

In [19]: viaduct_img

Out[19]:



```
In [20]: get_top_5_predictions(viaduct_img)
```

```
Out[20]: prediction
1      viaduct
2      pier
3      dam
4      prison
5  solar dish
```

Image 3: water tower

```
In [21]: water_tower_img
```

```
Out[21]:
```



```
In [22]: get_top_5_predictions(water_tower_img)
```

```
Out[22]:          prediction
1         solar dish
2         water tower
3  aircraft carrier
4     jigsaw puzzle
5         oxygen mask
```